An Introduction to BigQuery

(in less than 10 minutes)

brought to you by

The ISB Cancer Genomics Cloud

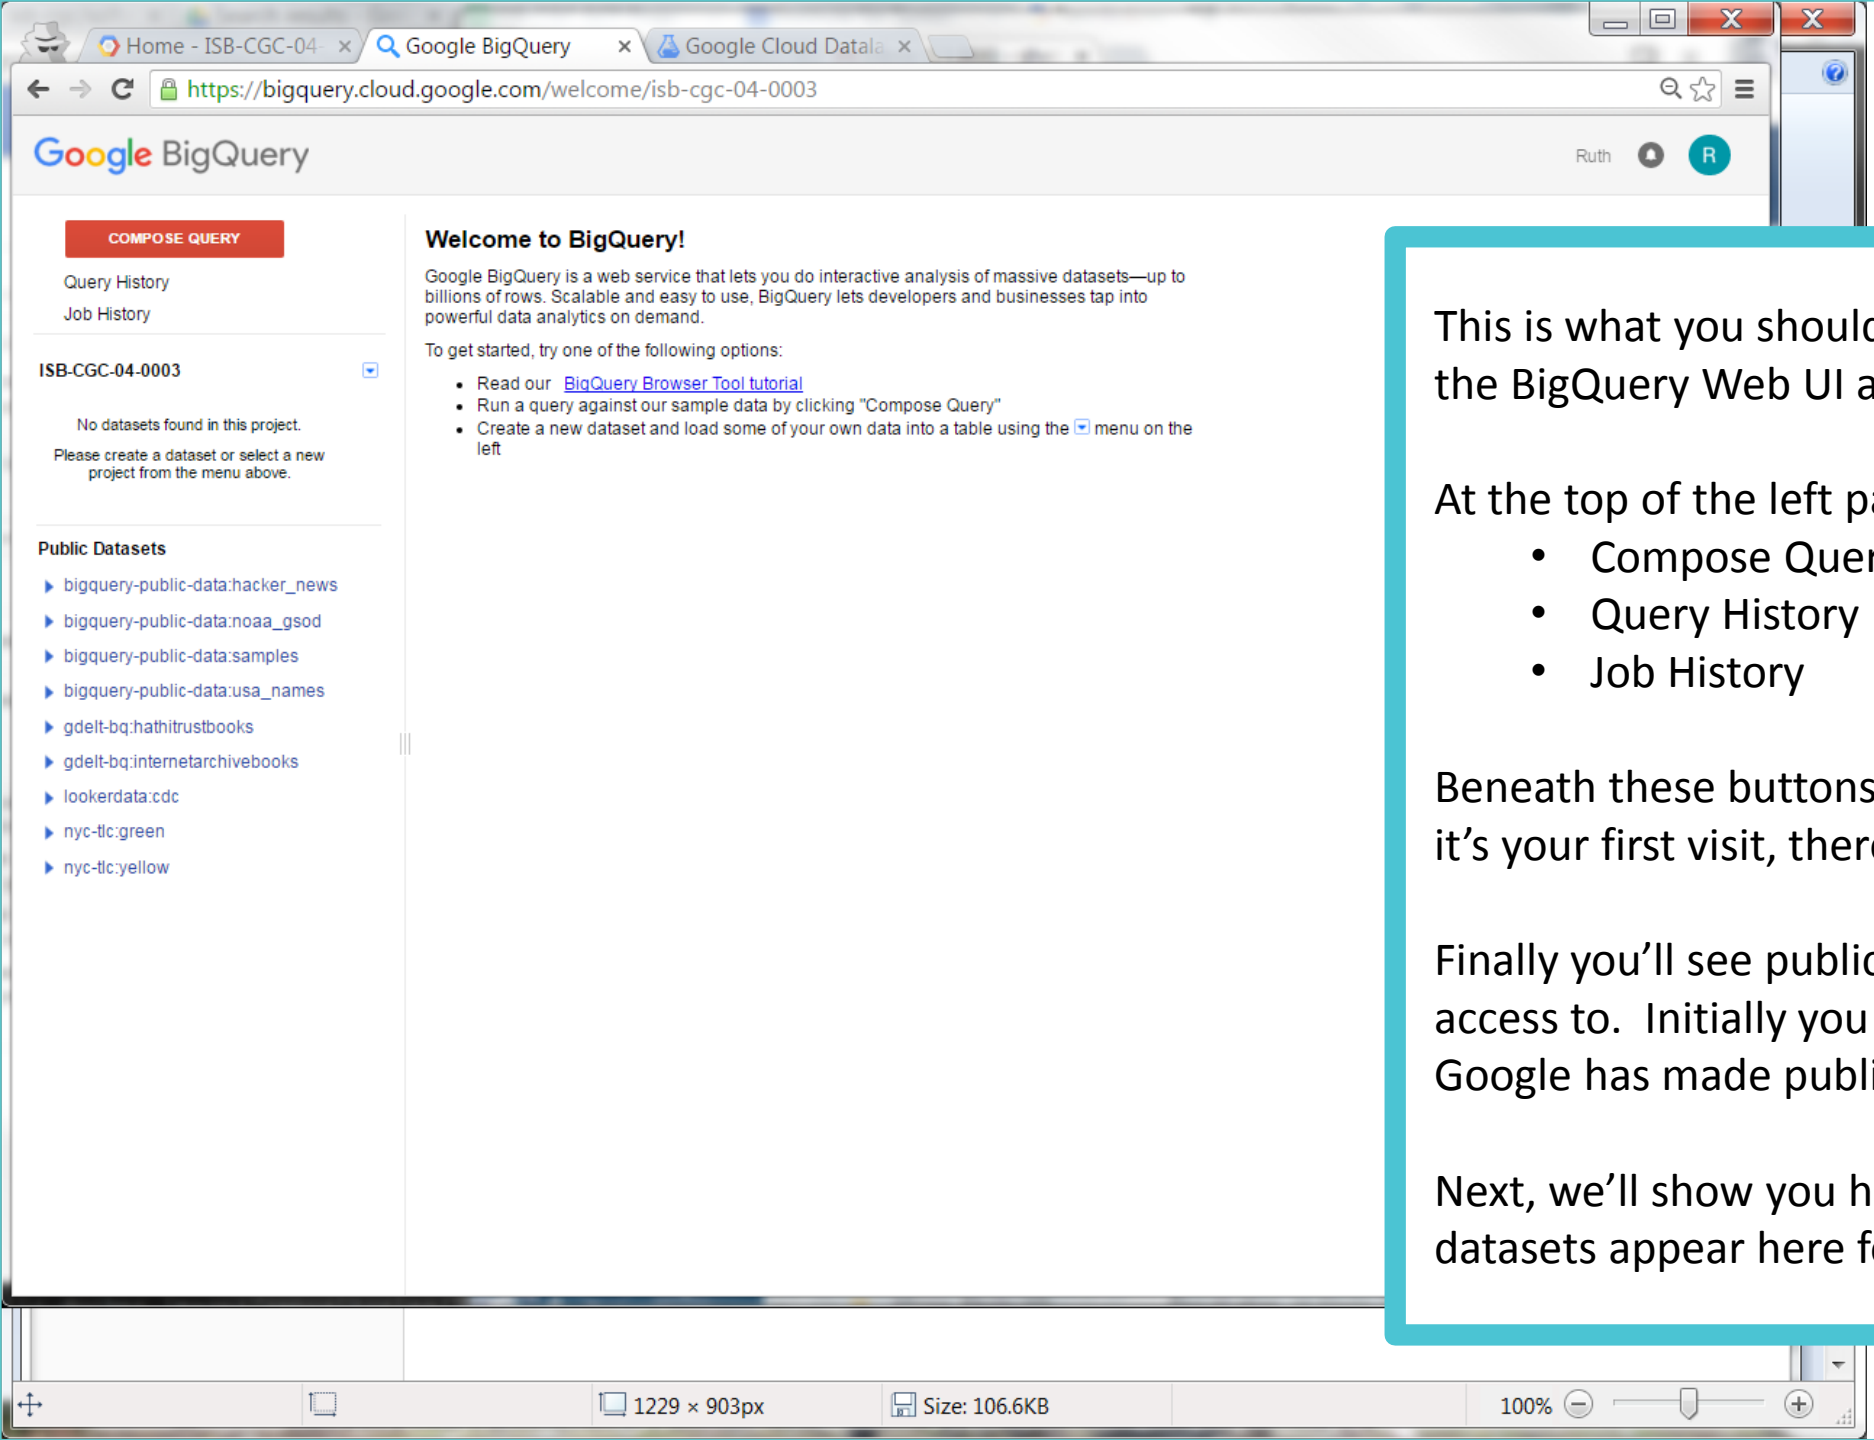

This is what you should see the first time you go to the BigQuery Web UI at bigquery.cloud.google.com

At the top of the left panel are three buttons:

- Compose Query
- Query History
- Job History

Beneath these buttons is your project space. Since it's your first visit, there are no datasets.

Finally you'll see public datasets that you may have access to. Initially you will see a few datasets that Google has made public.

Next, we'll show you how to make the ISB-CGC datasets appear here for easy access.

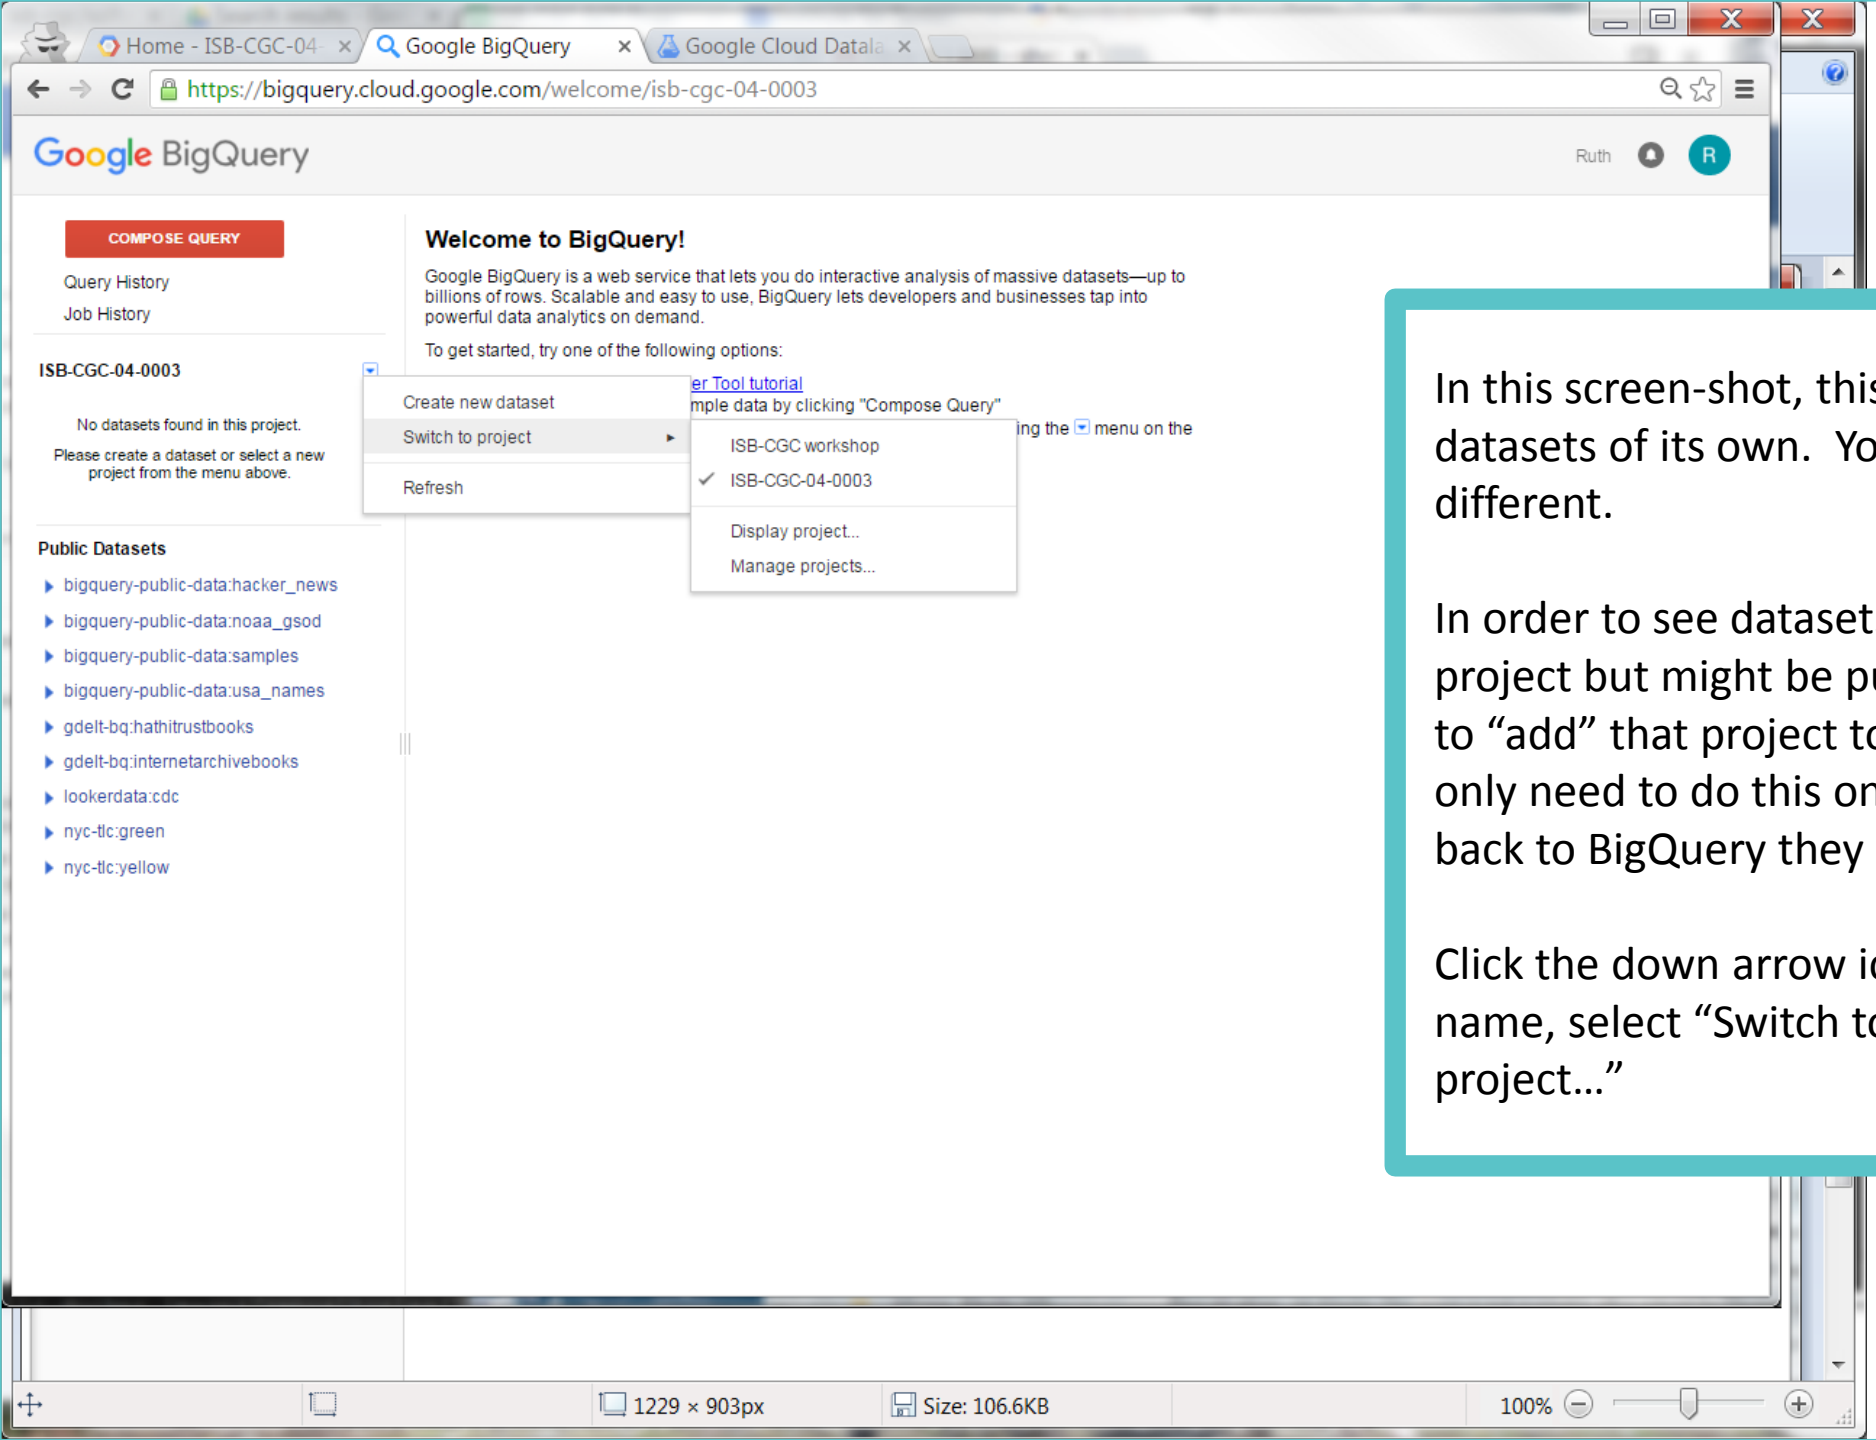

In this screen-shot, this particular project has no datasets of its own. Your project might look different.

In order to see datasets that are owned by another project but might be publicly-accessible, you need to "add" that project to your BigQuery view. (You'll only need to do this once – next time you come back to BigQuery they will already be there.)

Click the down arrow icon next to your project name, select "Switch to project", and then "Display project…"

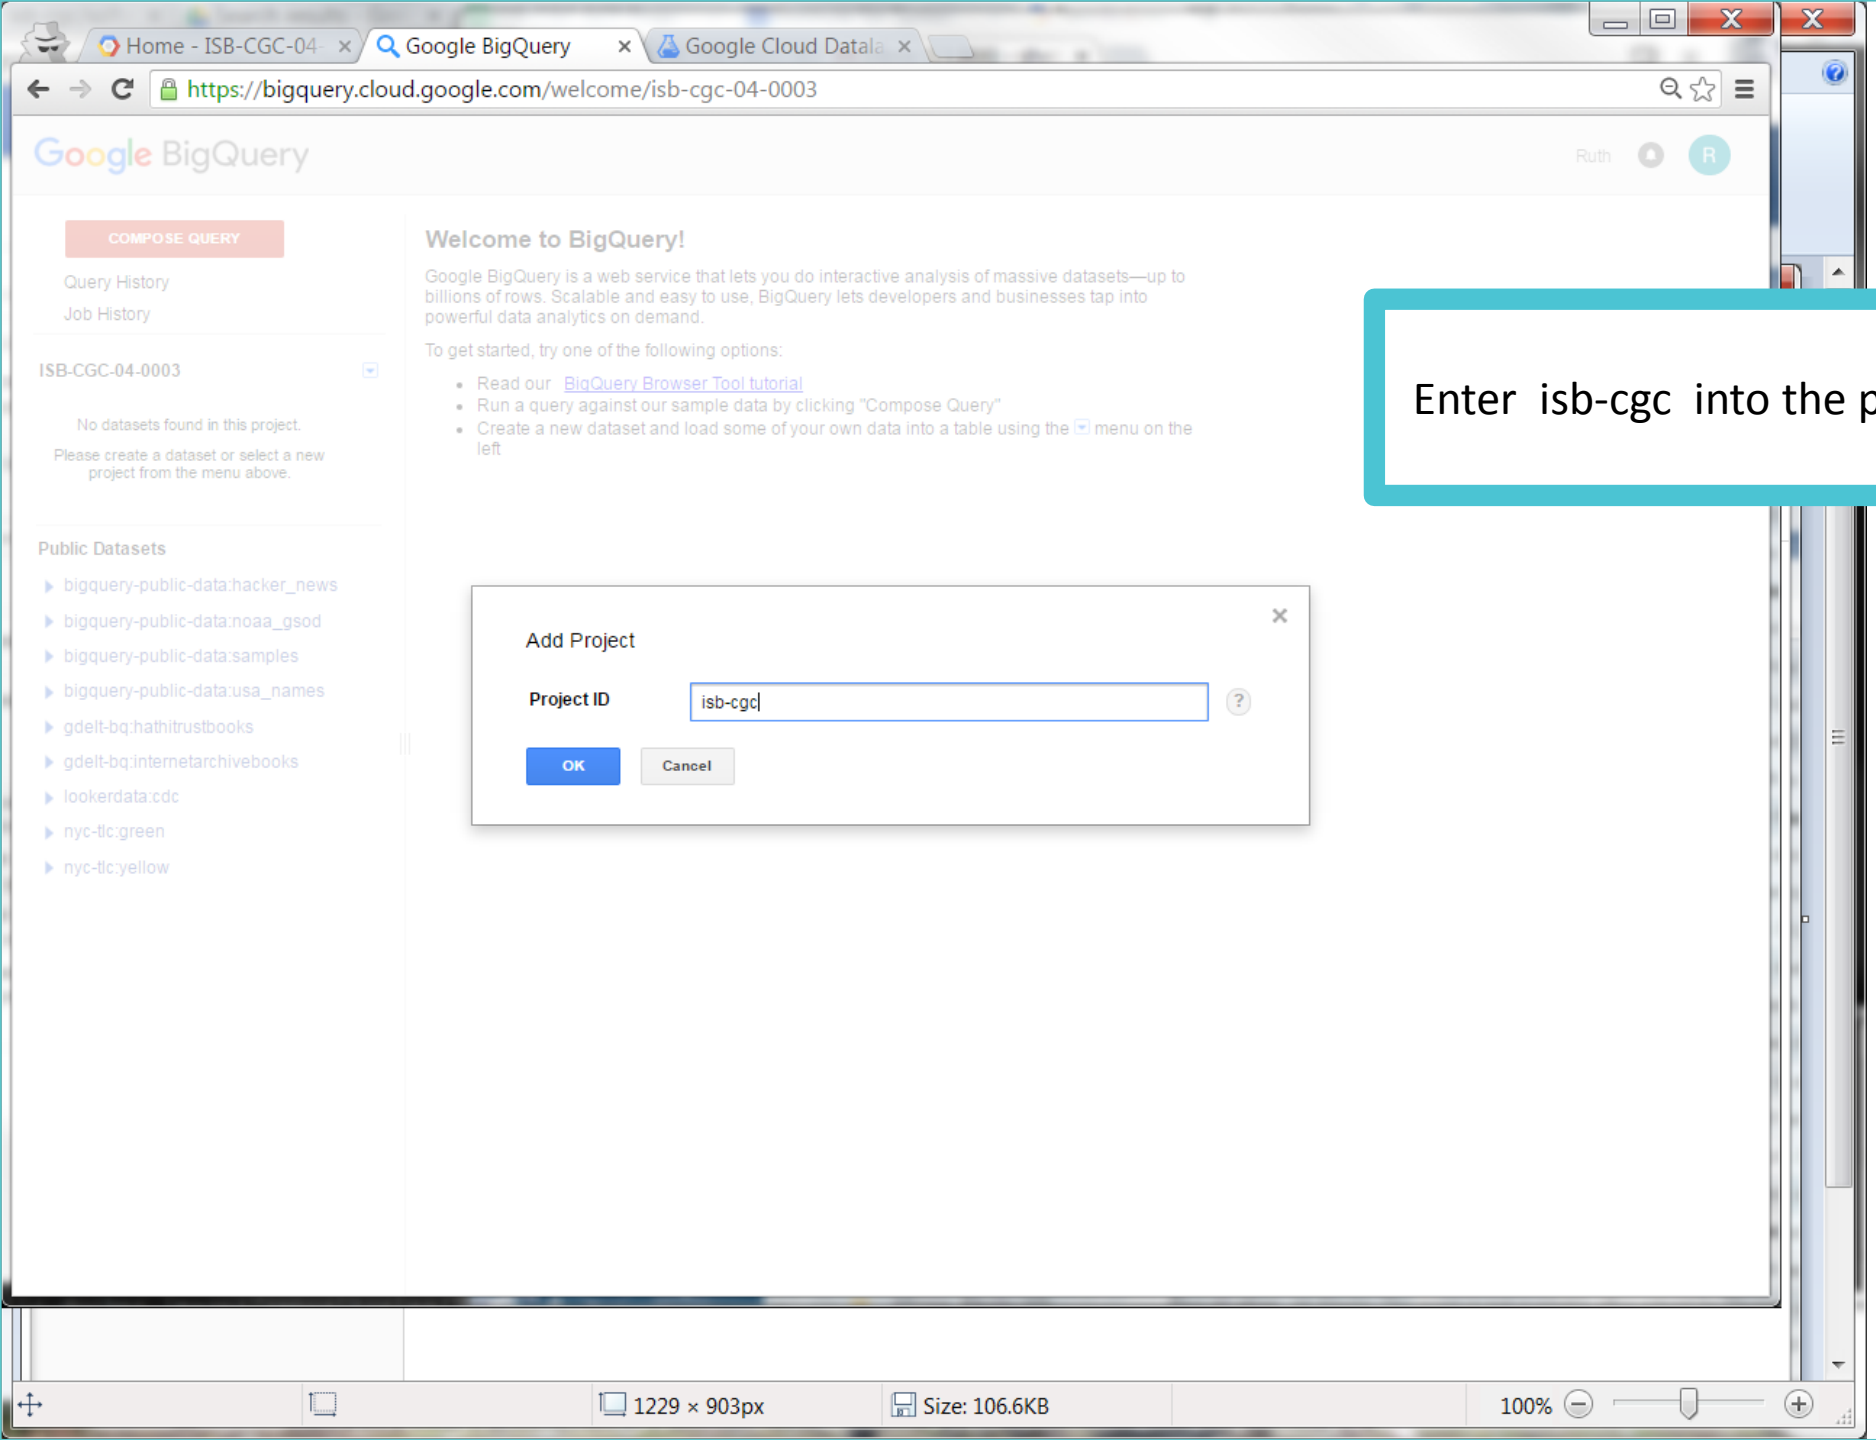

Enter isb-cgc into the pop-up window, and click **OK**

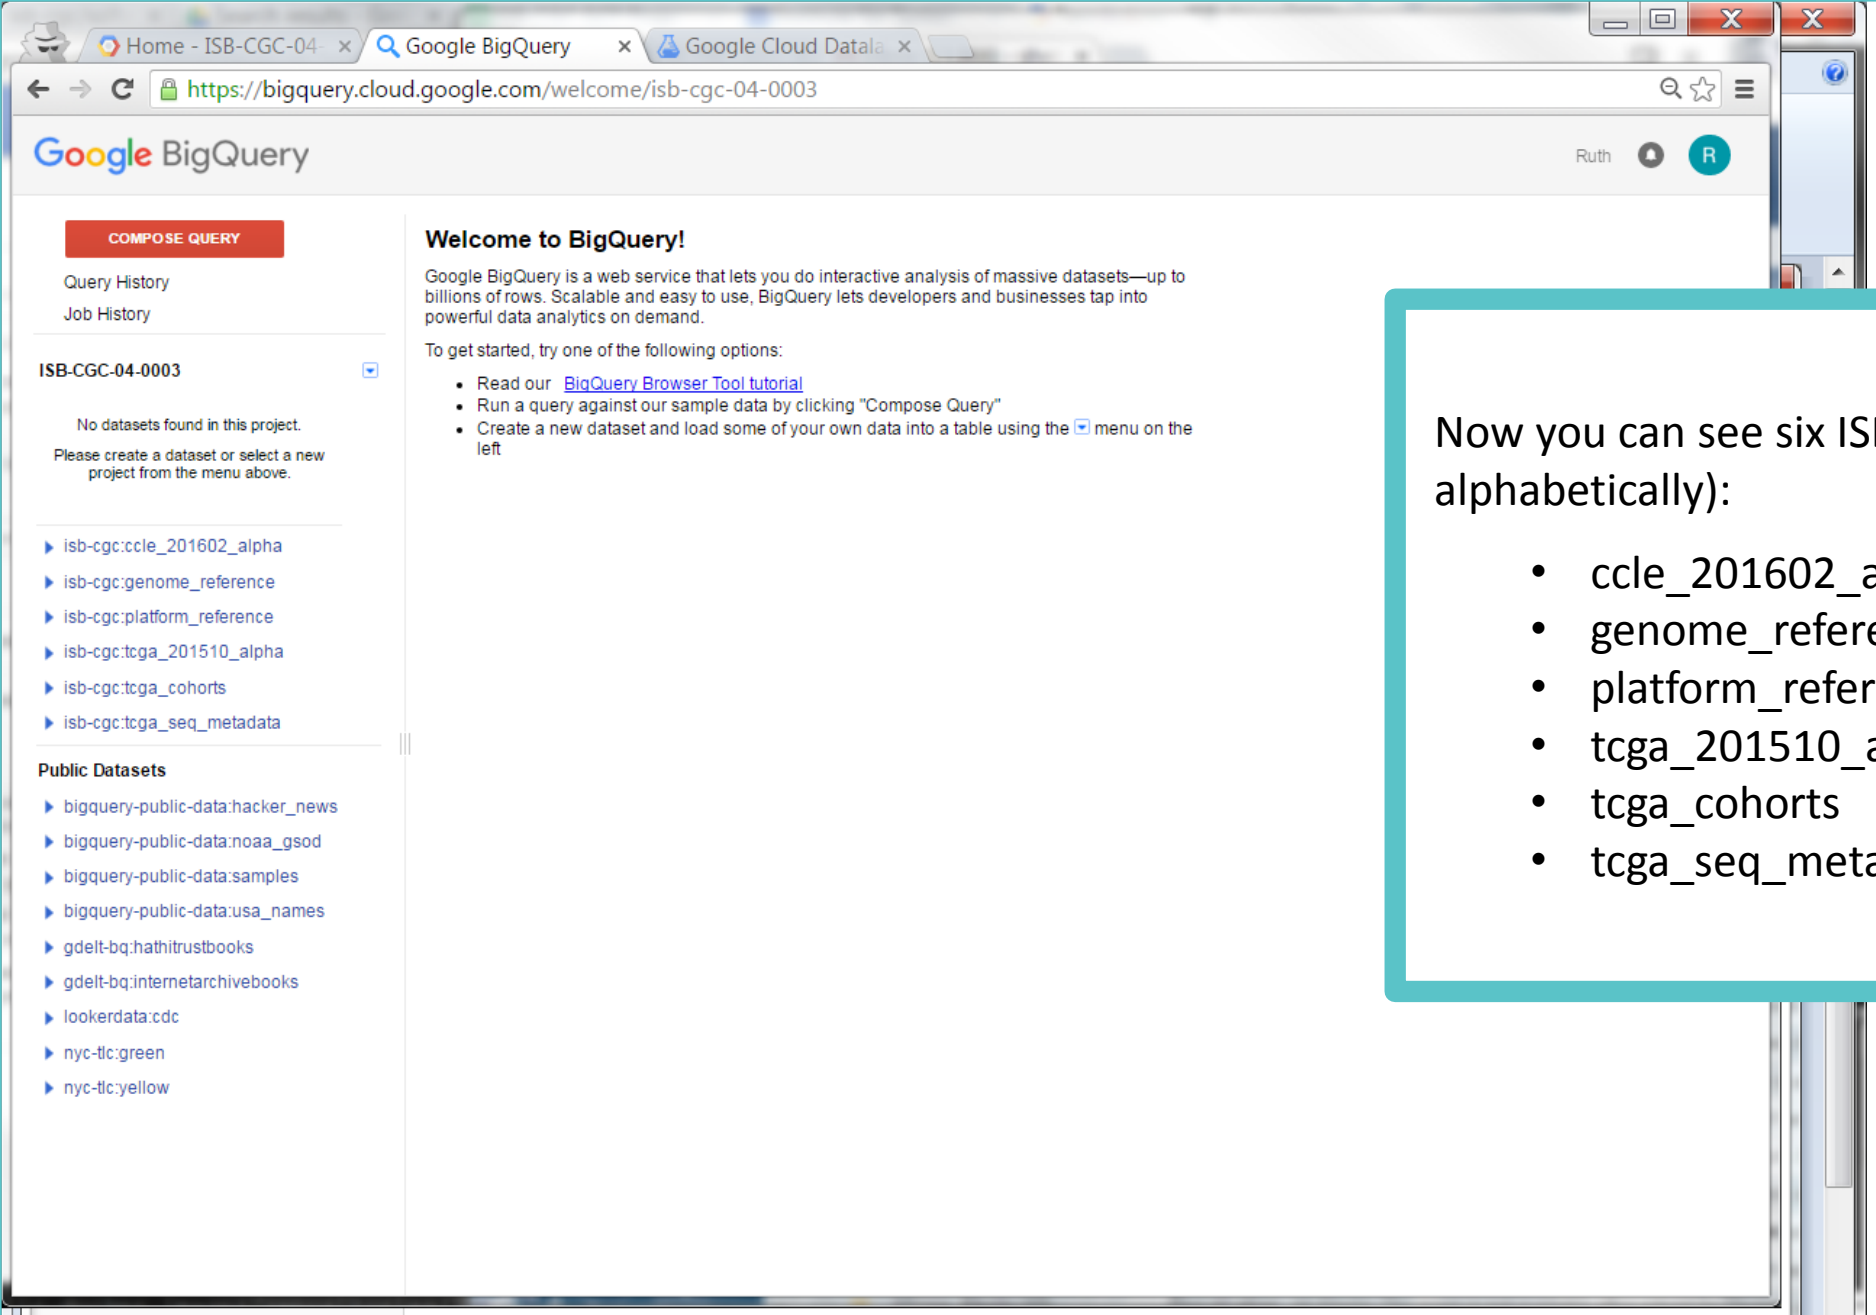

Now you can see six ISB-CGC datasets (arranged alphabetically):

- ccle_201602_alpha
- genome_reference
- platform_reference
- tcga_201510_alpha
- tcga_cohorts
- tcga_seq_metadata

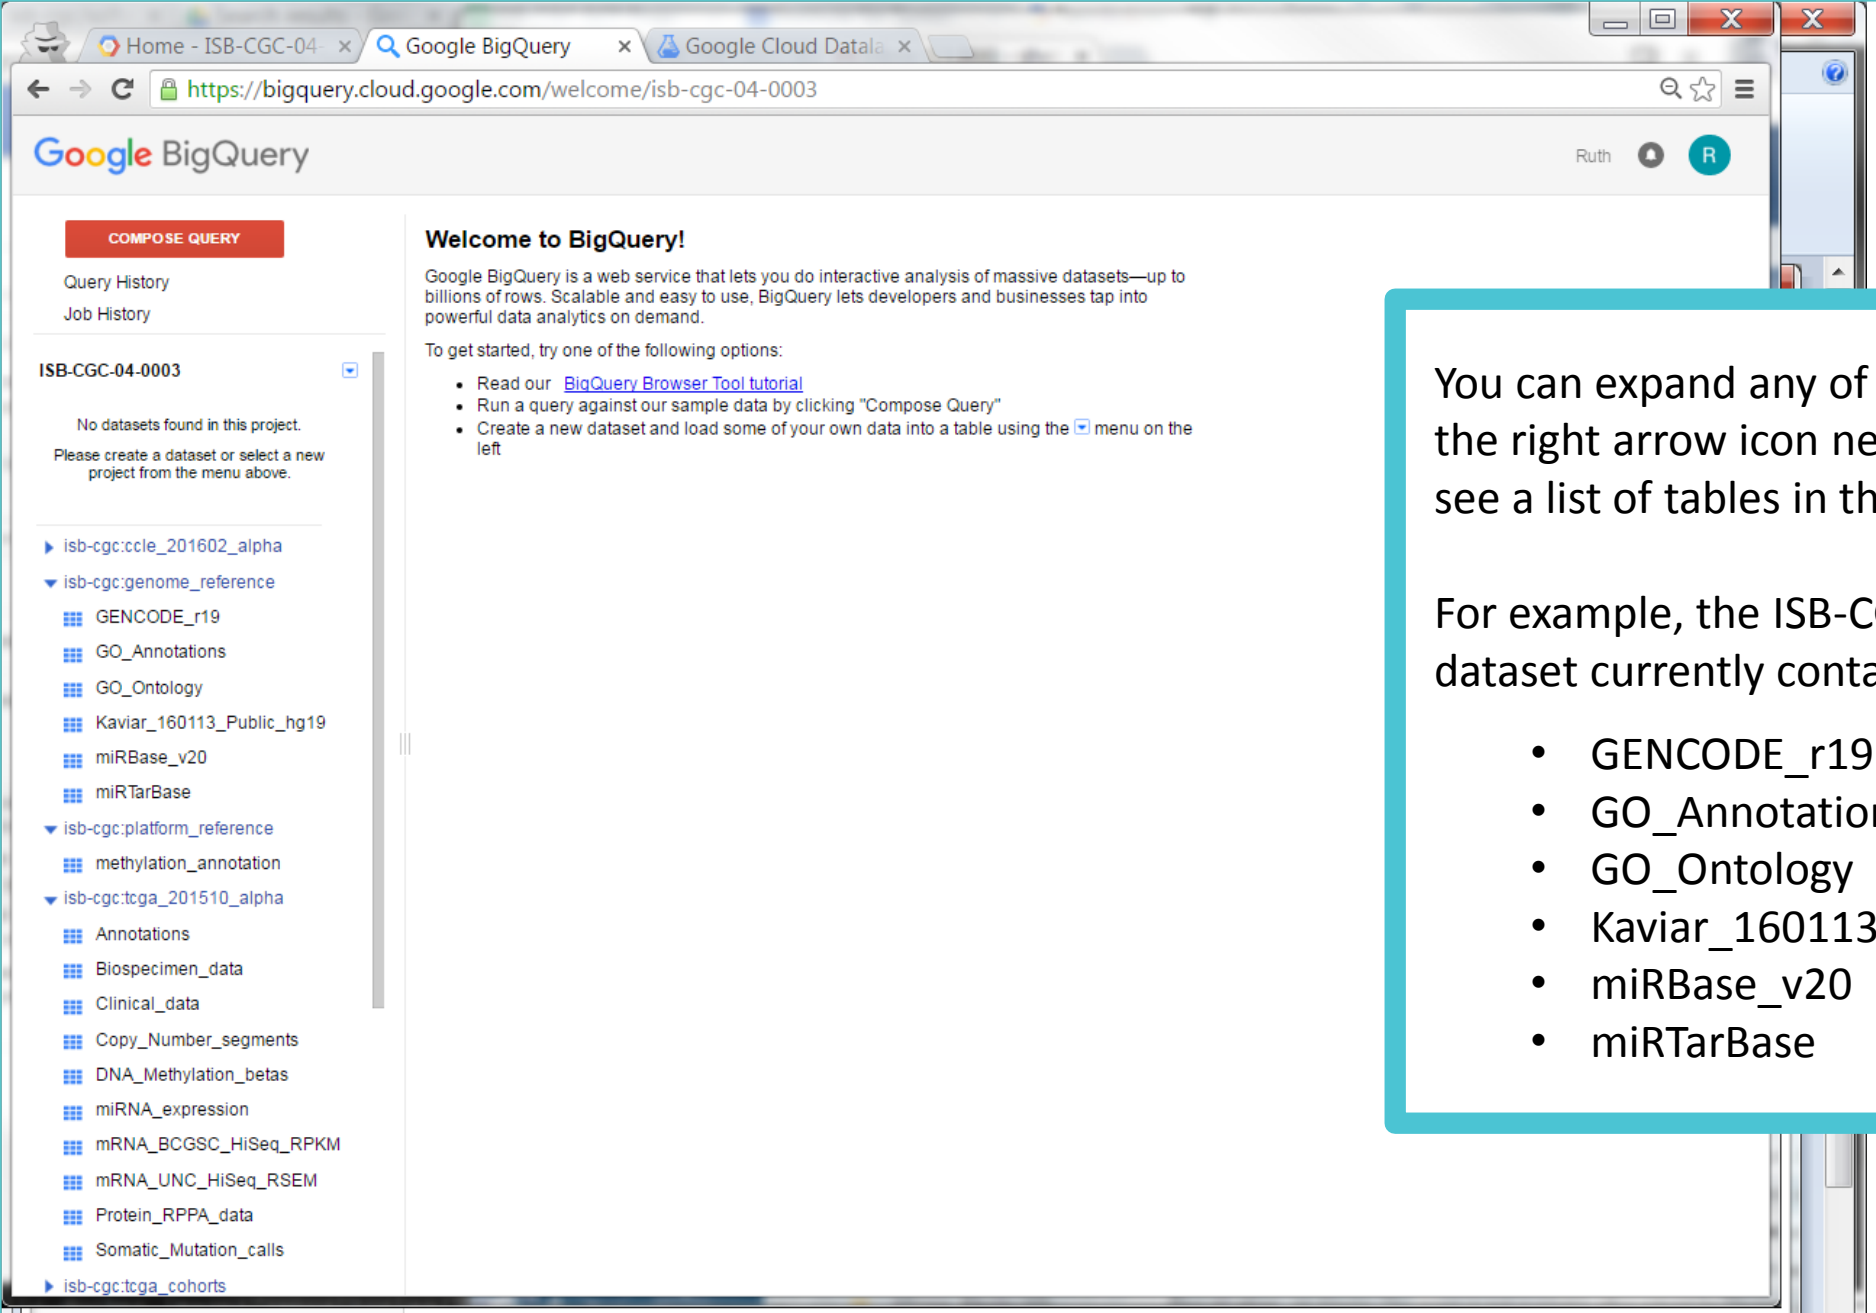

You can expand any of the datasets by clicking on the right arrow icon next to the dataset name, to see a list of tables in that dataset.

For example, the ISB-CGC “genome_reference” dataset currently contains the following tables:

- GENCODE_r19
- GO_Annotations
- GO_Ontology
- Kaviar_160113_Public_hg19
- miRBase_v20
- miRTarBase

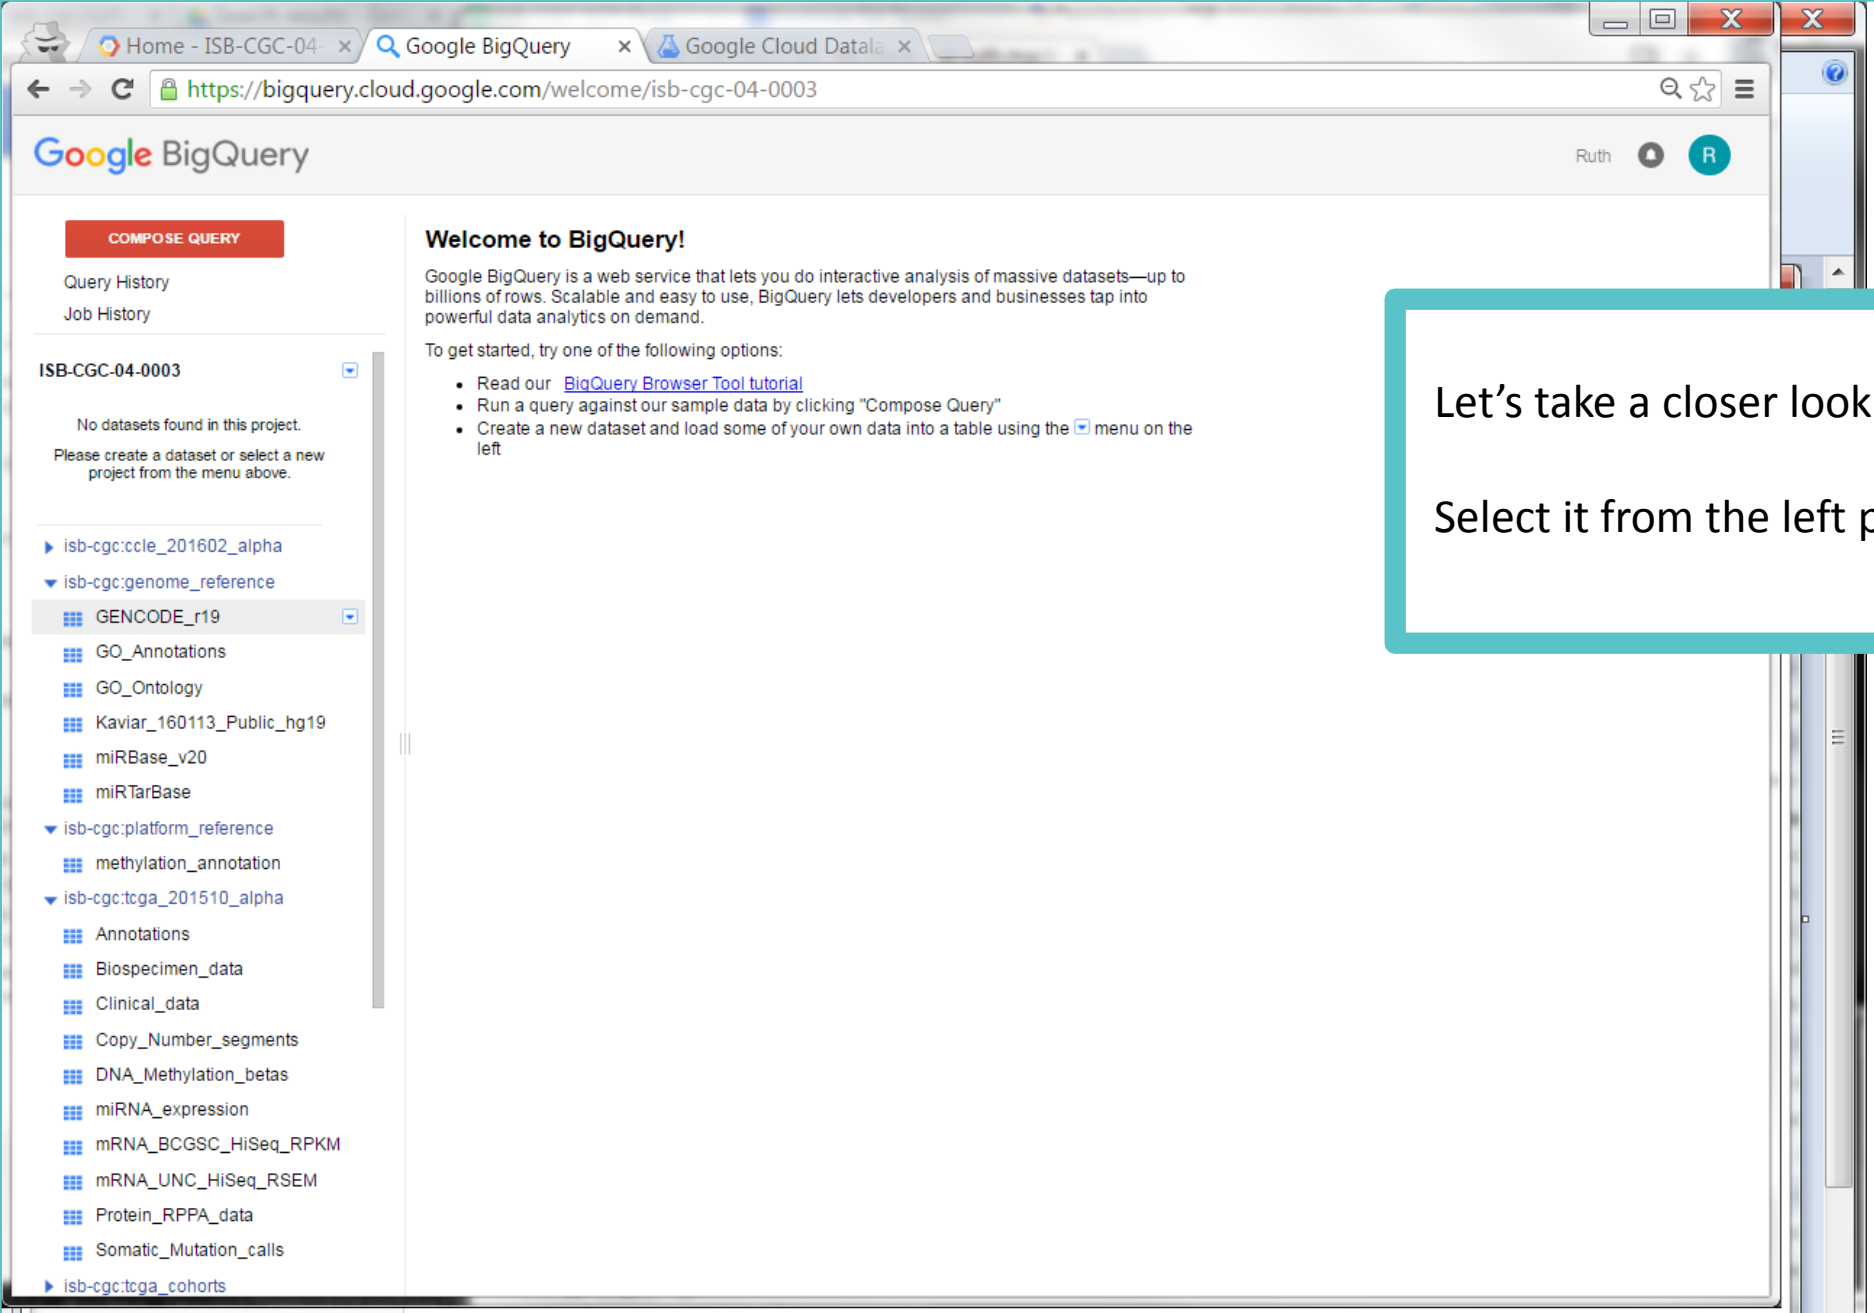

Let's take a closer look at the GENCODE_r19 table.

Select it from the left panel and click.

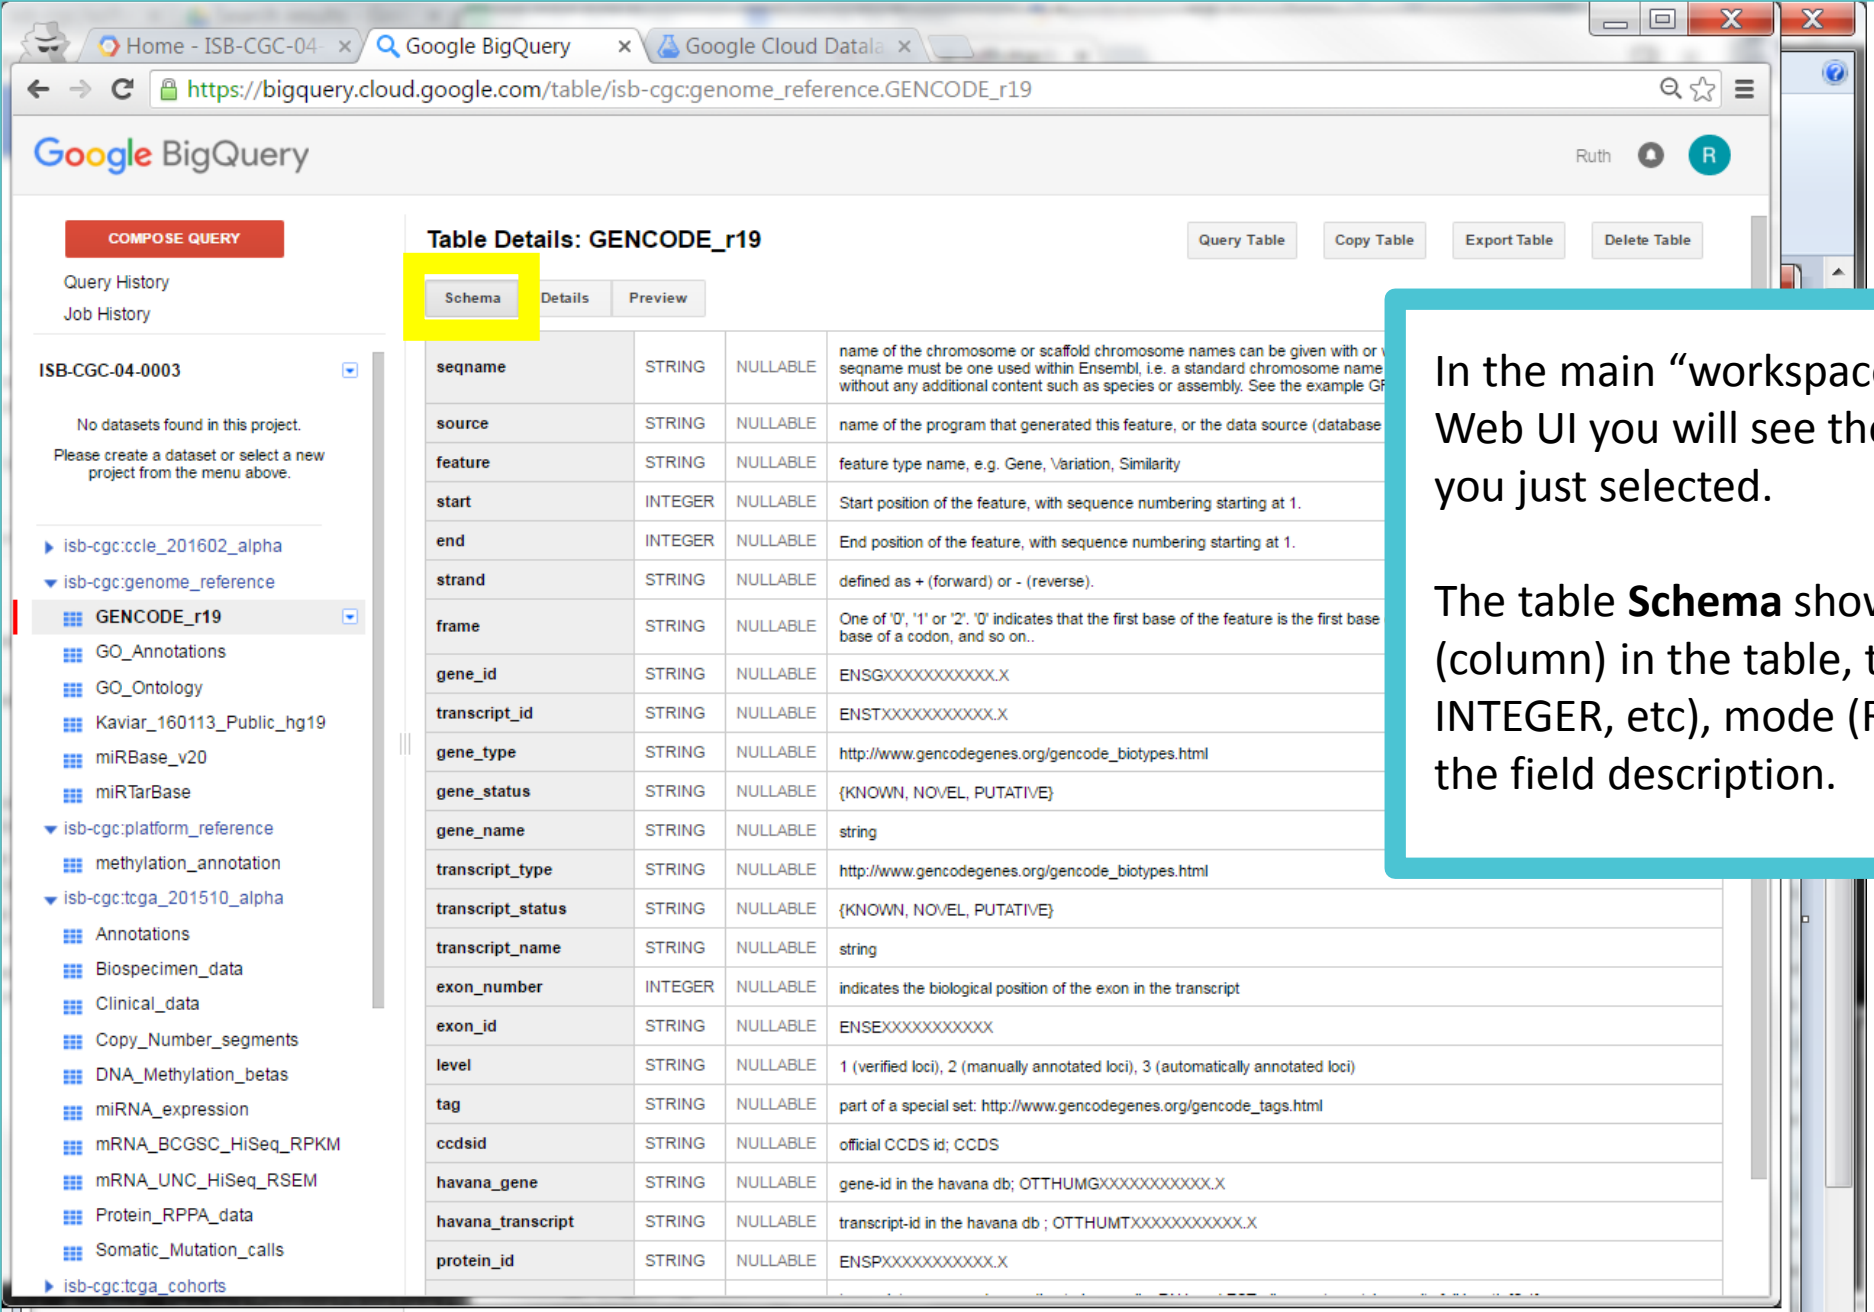

In the main “workspace” portion of the BigQuery Web UI you will see the “Table Details” for the table you just selected.

The table **Schema** shows the name of each field (column) in the table, the data type (STRING, INTEGER, etc), mode (REQUIRED or NULLABLE), and the field description.

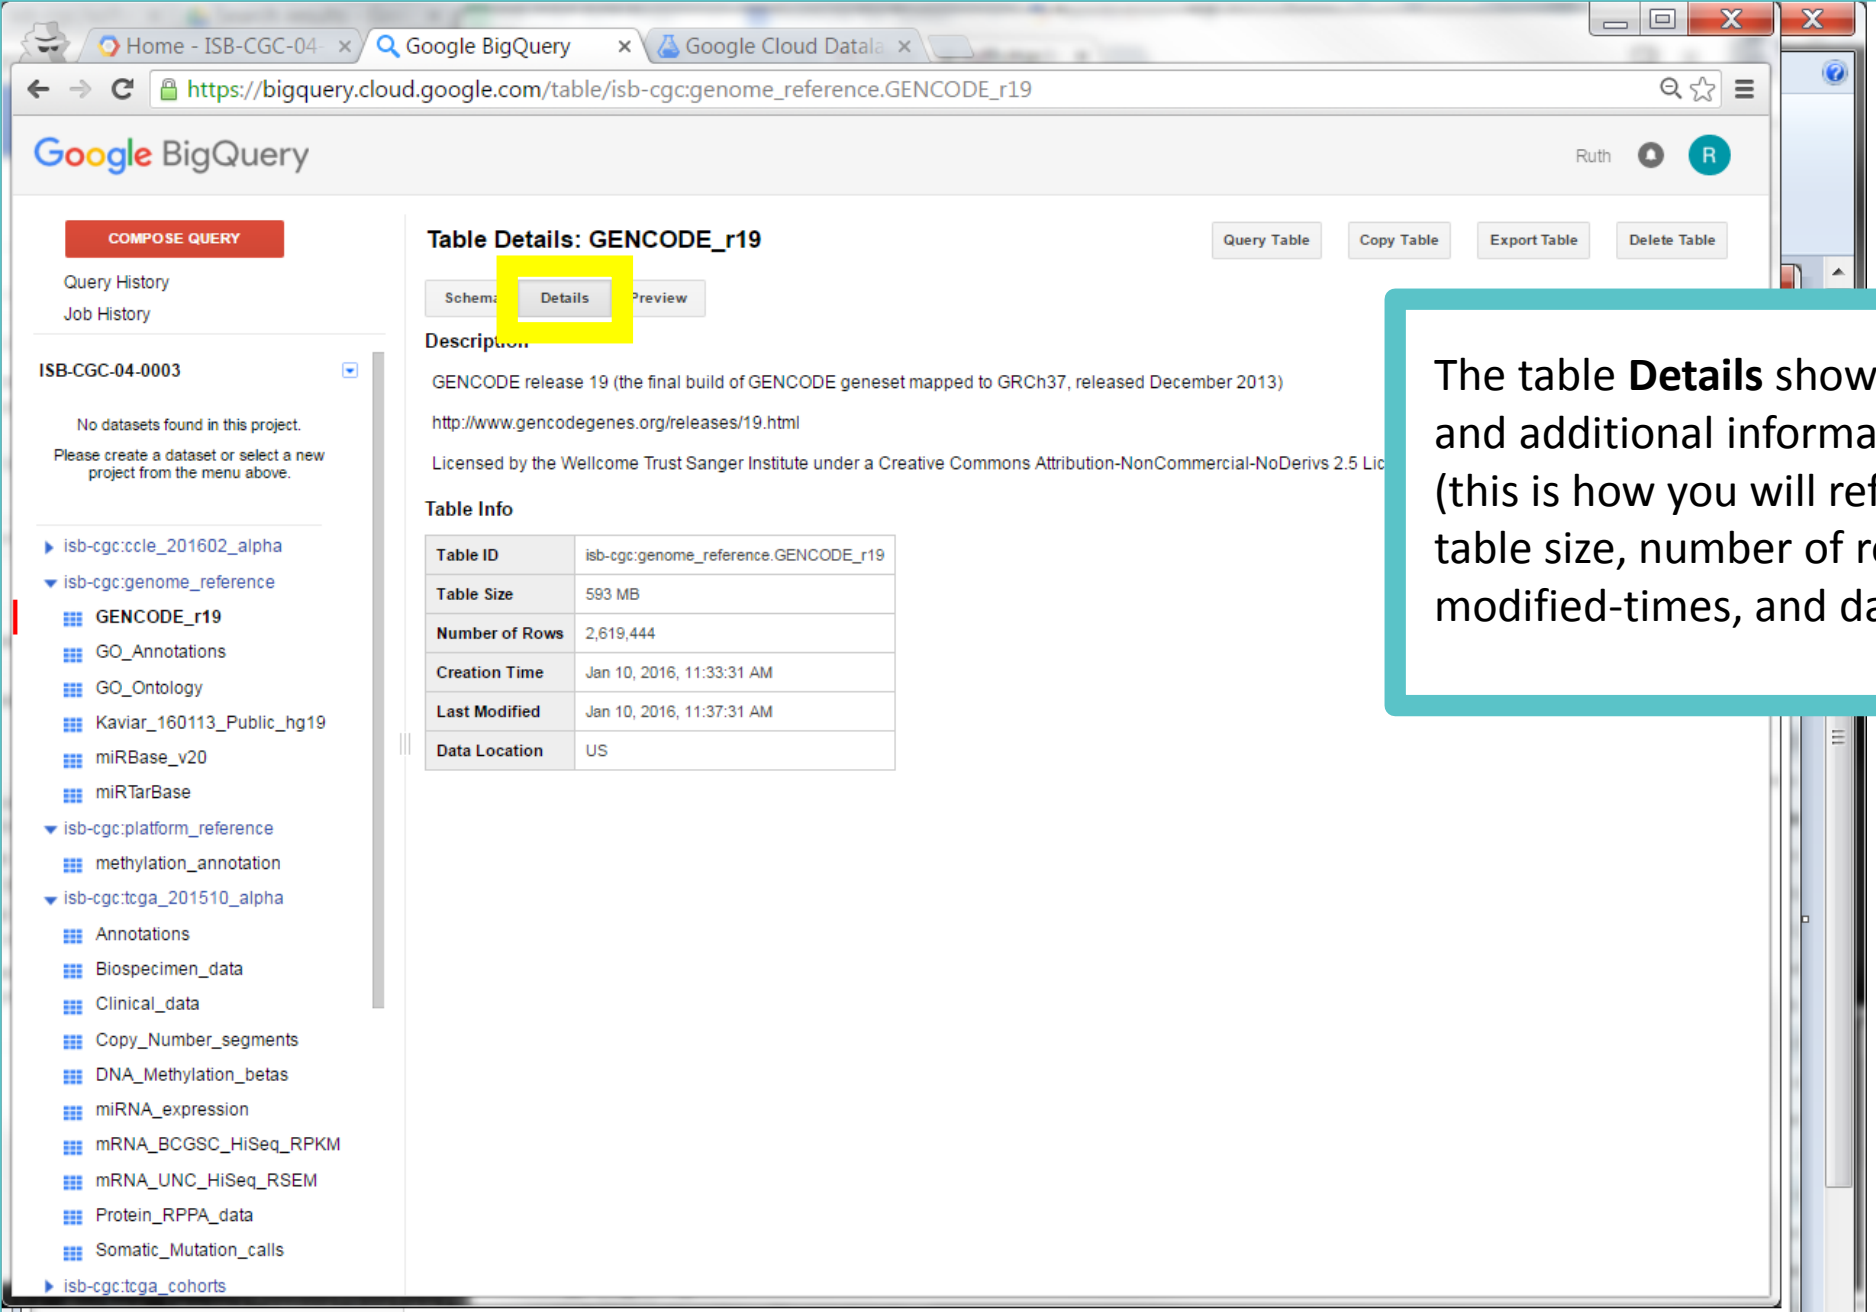

The table **Details** shows you the table Description and additional information including the table ID (this is how you will refer to it in a SQL query), the table size, number of rows, creation- and last-modified-times, and data location.

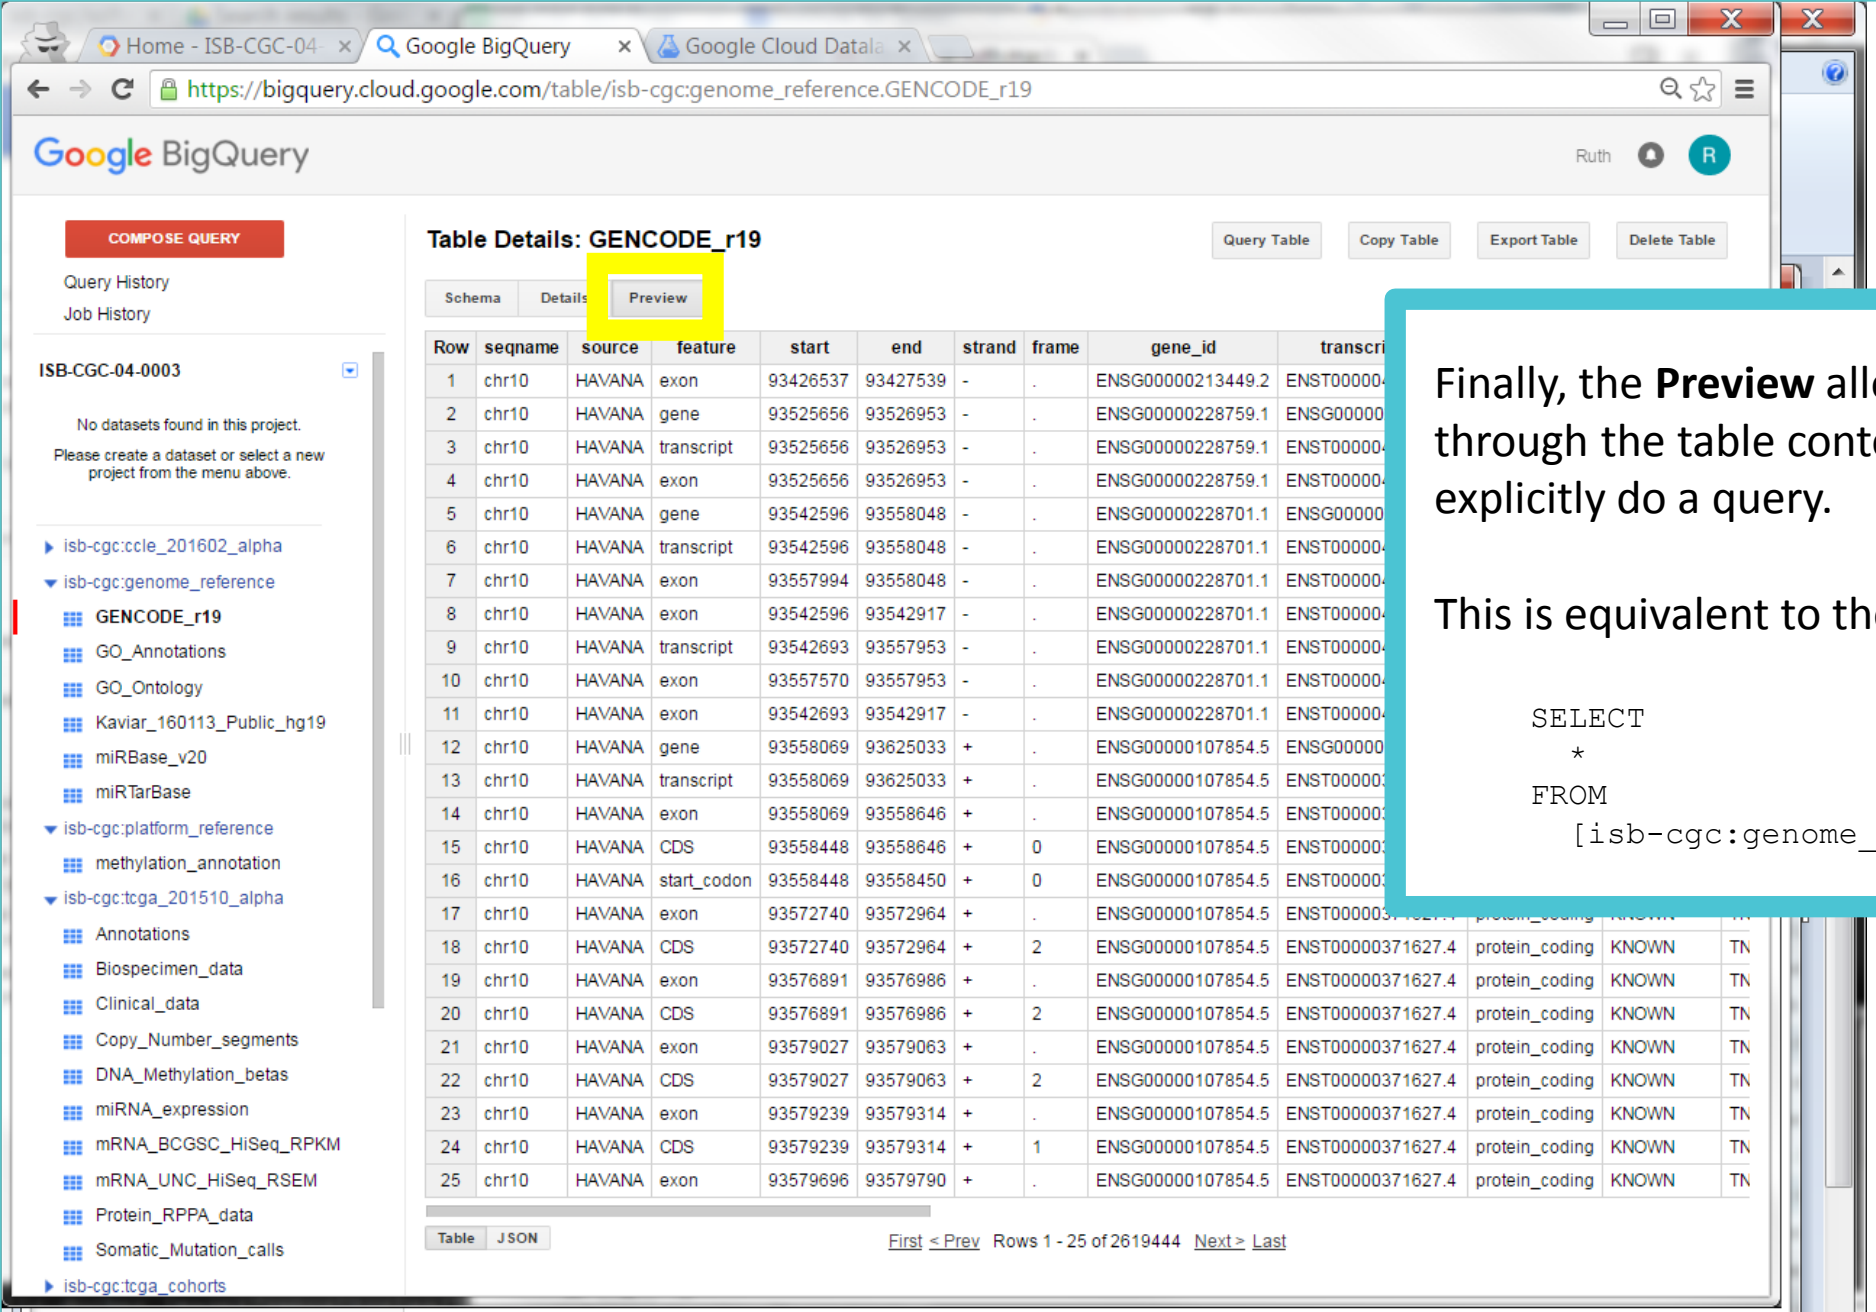

Finally, the **Preview** allows you see to see and scroll through the table contents without having to explicitly do a query.

This is equivalent to the following SQL query:

```
SELECT
  *
FROM
  [isb-cgc:genome_reference.GENCODE_r19]
```

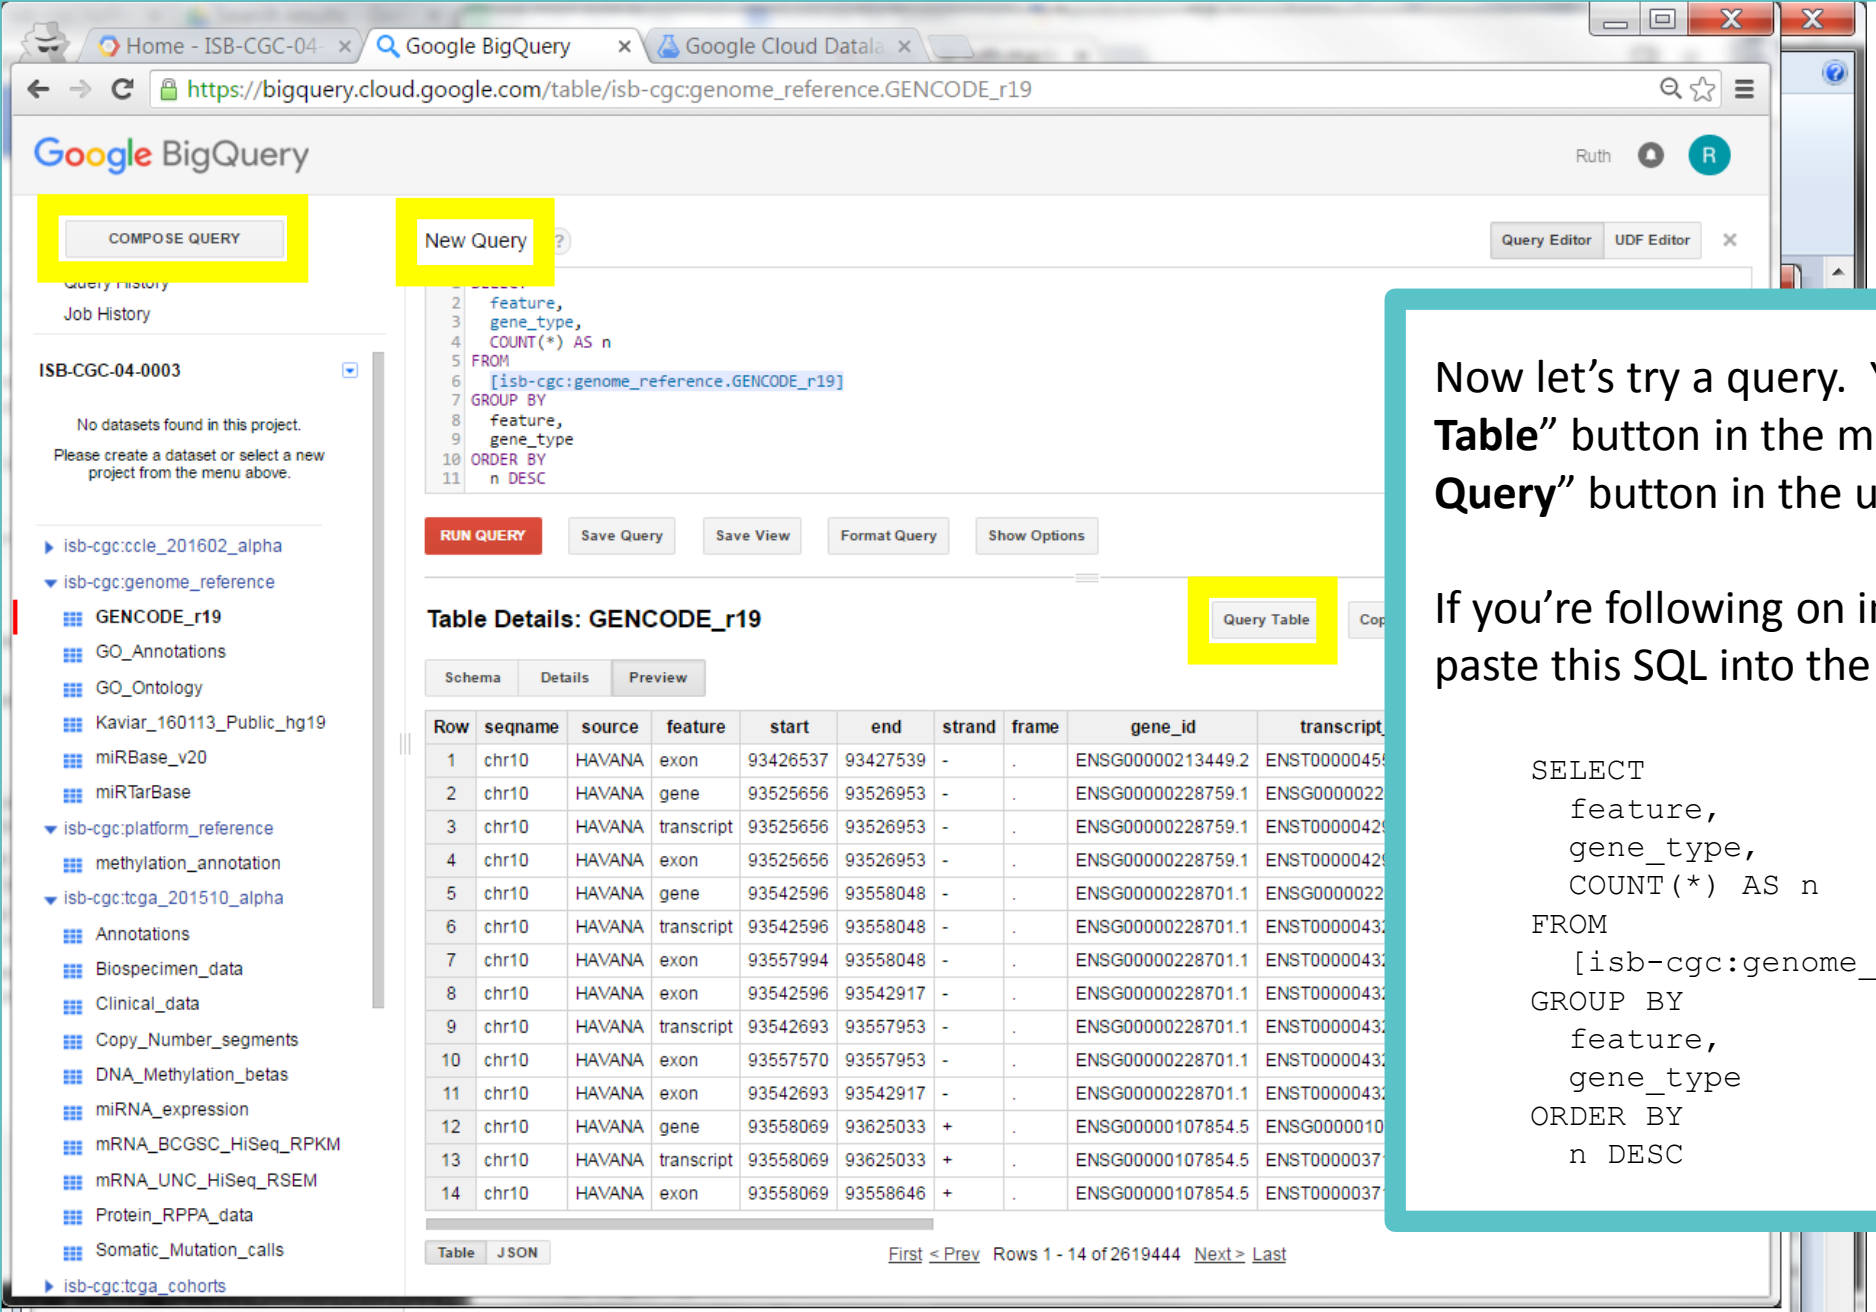

Now let's try a query. You can click on the "**Query Table**" button in the main panel or in the "**Compose Query**" button in the upper left corner.

If you're following on in your own browser, cut and paste this SQL into the **New Query** text area:

```
SELECT
  feature,
  gene_type,
  COUNT(*) AS n
FROM
  [isb-cgc:genome_reference.GENCODE_r19]
GROUP BY
  feature,
  gene_type
ORDER BY
  n DESC
```

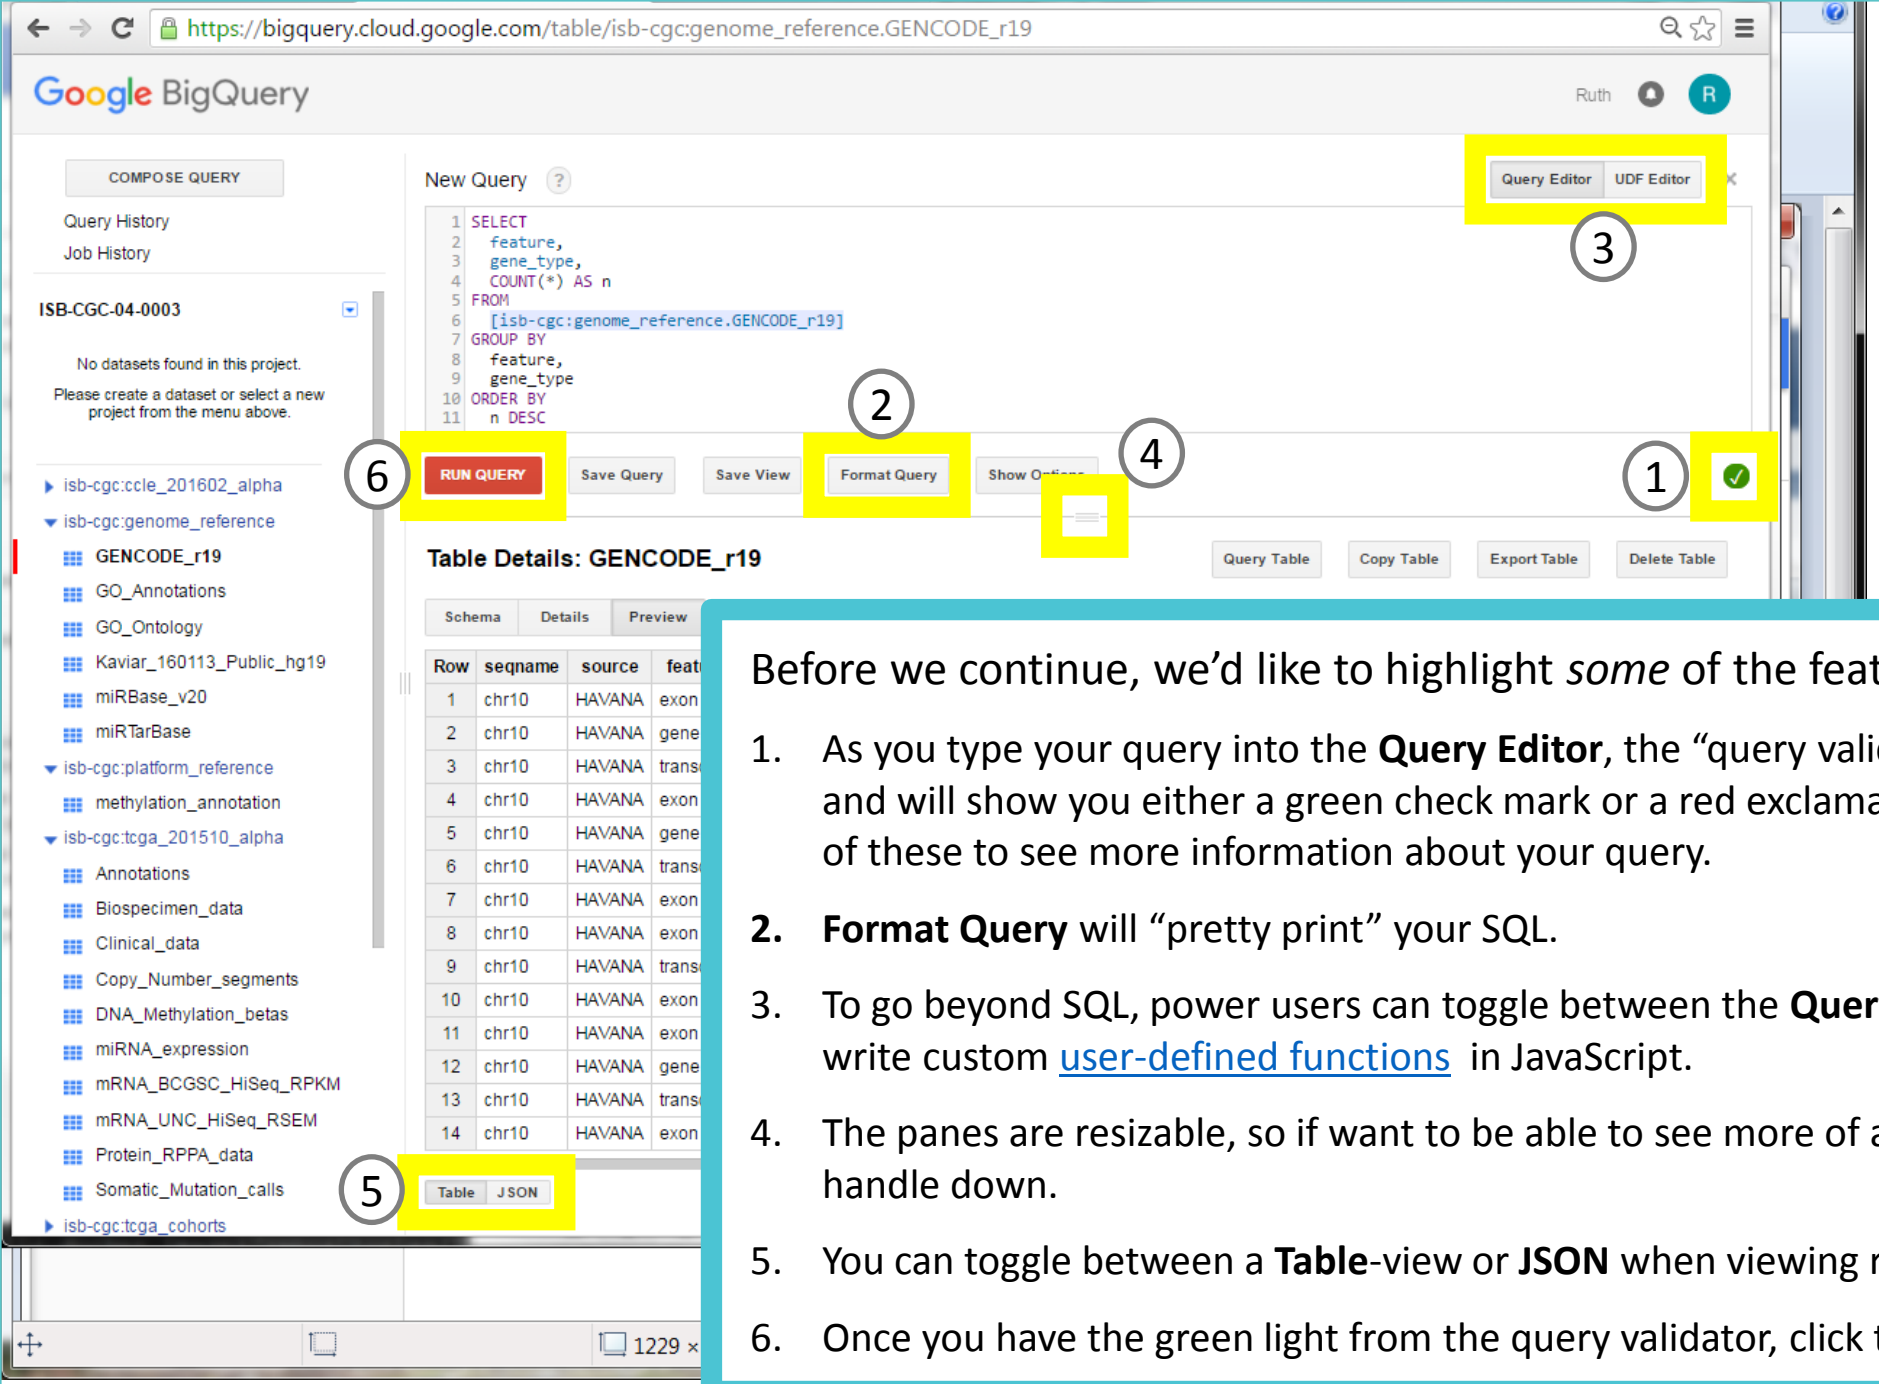

Before we continue, we'd like to highlight *some* of the features in the BigQuery Web UI:

1. As you type your query into the **Query Editor**, the "query validator" is automatically running, and will show you either a green check mark or a red exclamation point. You can click on either of these to see more information about your query.
2. **Format Query** will "pretty print" your SQL.
3. To go beyond SQL, power users can toggle between the **Query Editor** and the **UDF Editor** and write custom user-defined functions in JavaScript.
4. The panes are resizable, so if want to be able to see more of a long query you can drag the sash handle down.
5. You can toggle between a **Table**-view or **JSON** when viewing results.
6. Once you have the green light from the query validator, click the red **Run Query** button.

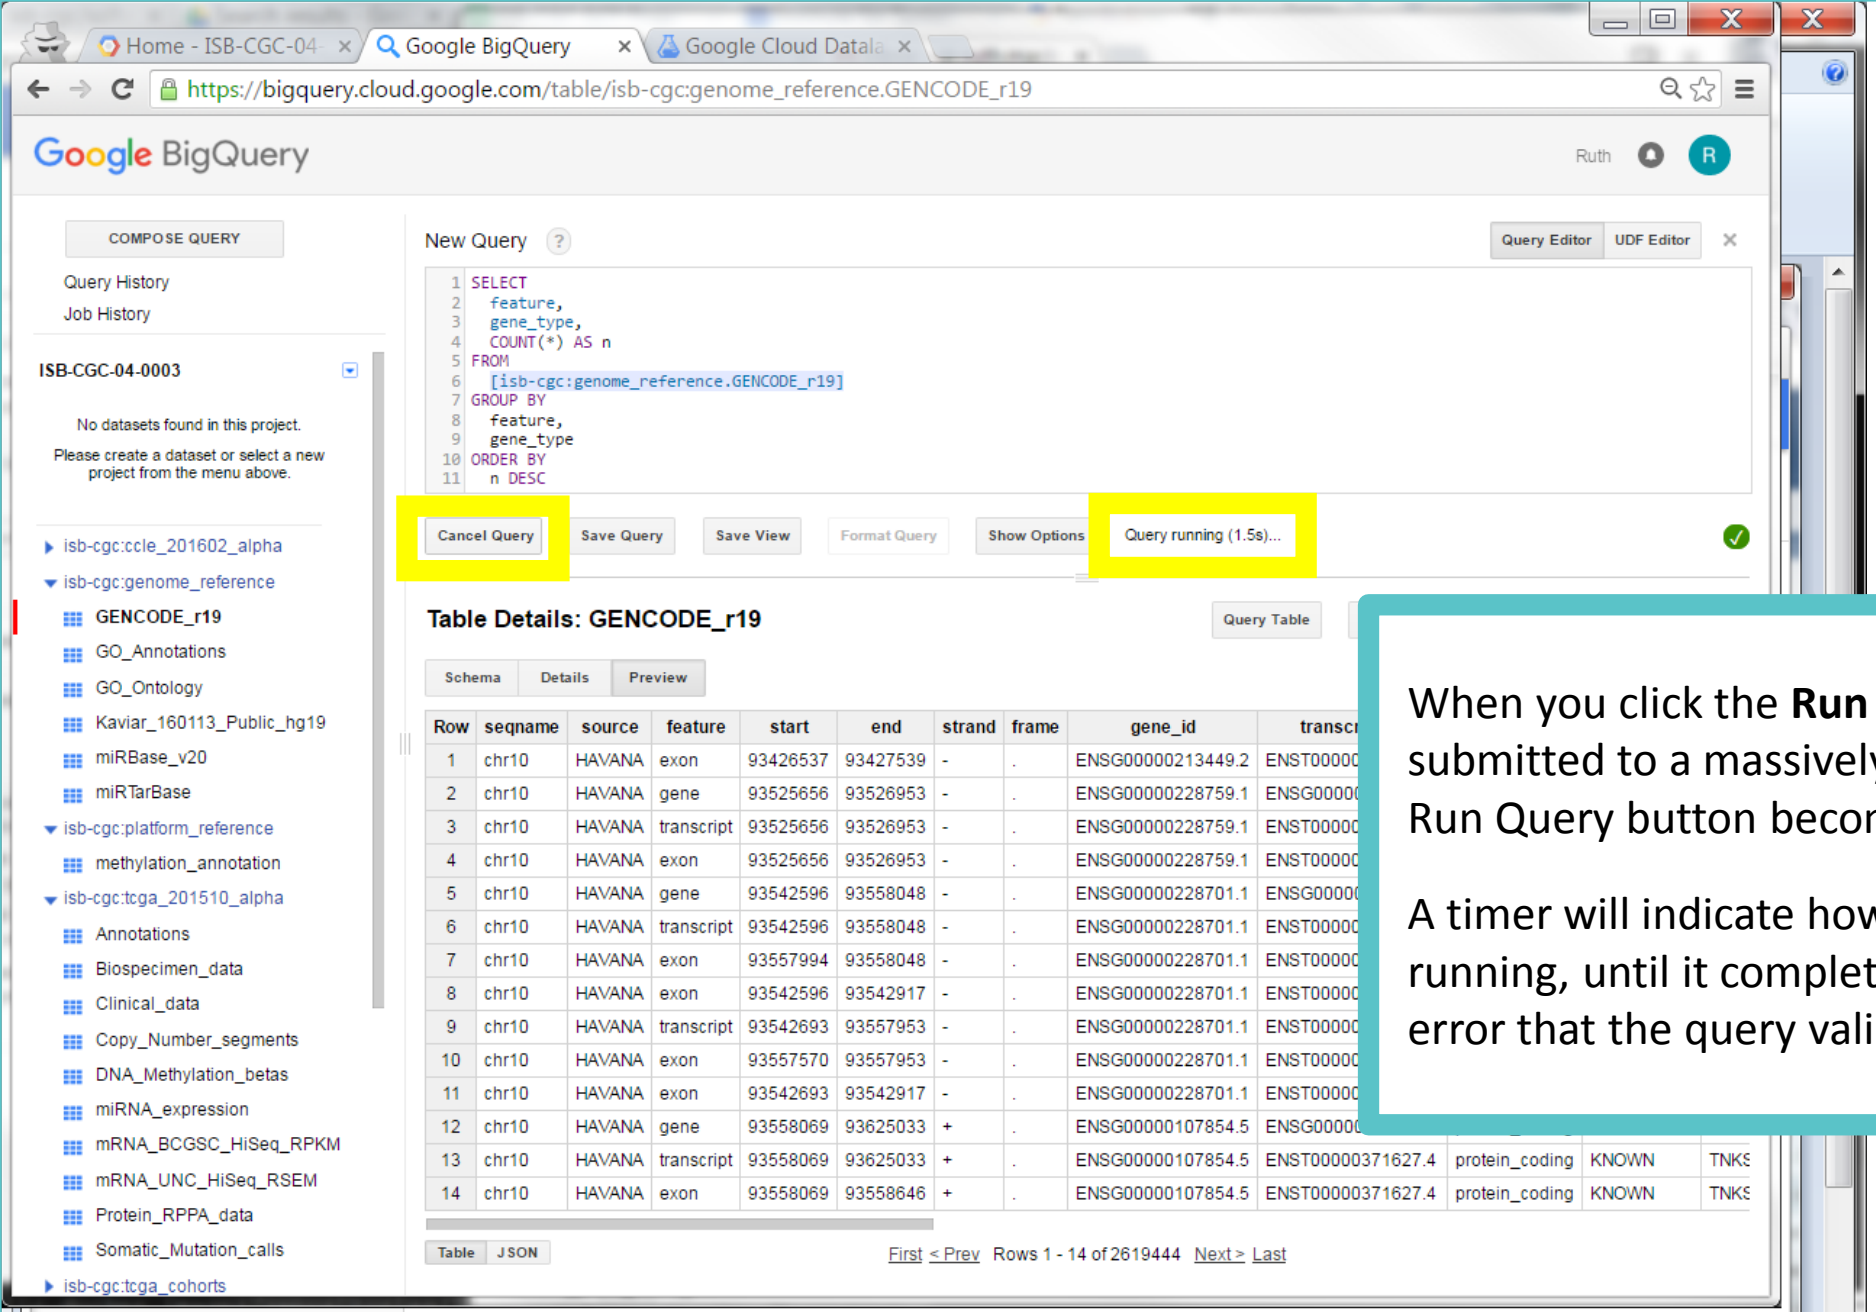

When you click the **Run Query** button, your query is submitted to a massively parallel engine (and the Run Query button becomes a **Cancel Query** button.)

A timer will indicate how long the query has been running, until it completes (or until it encounters an error that the query validator was not able to catch).

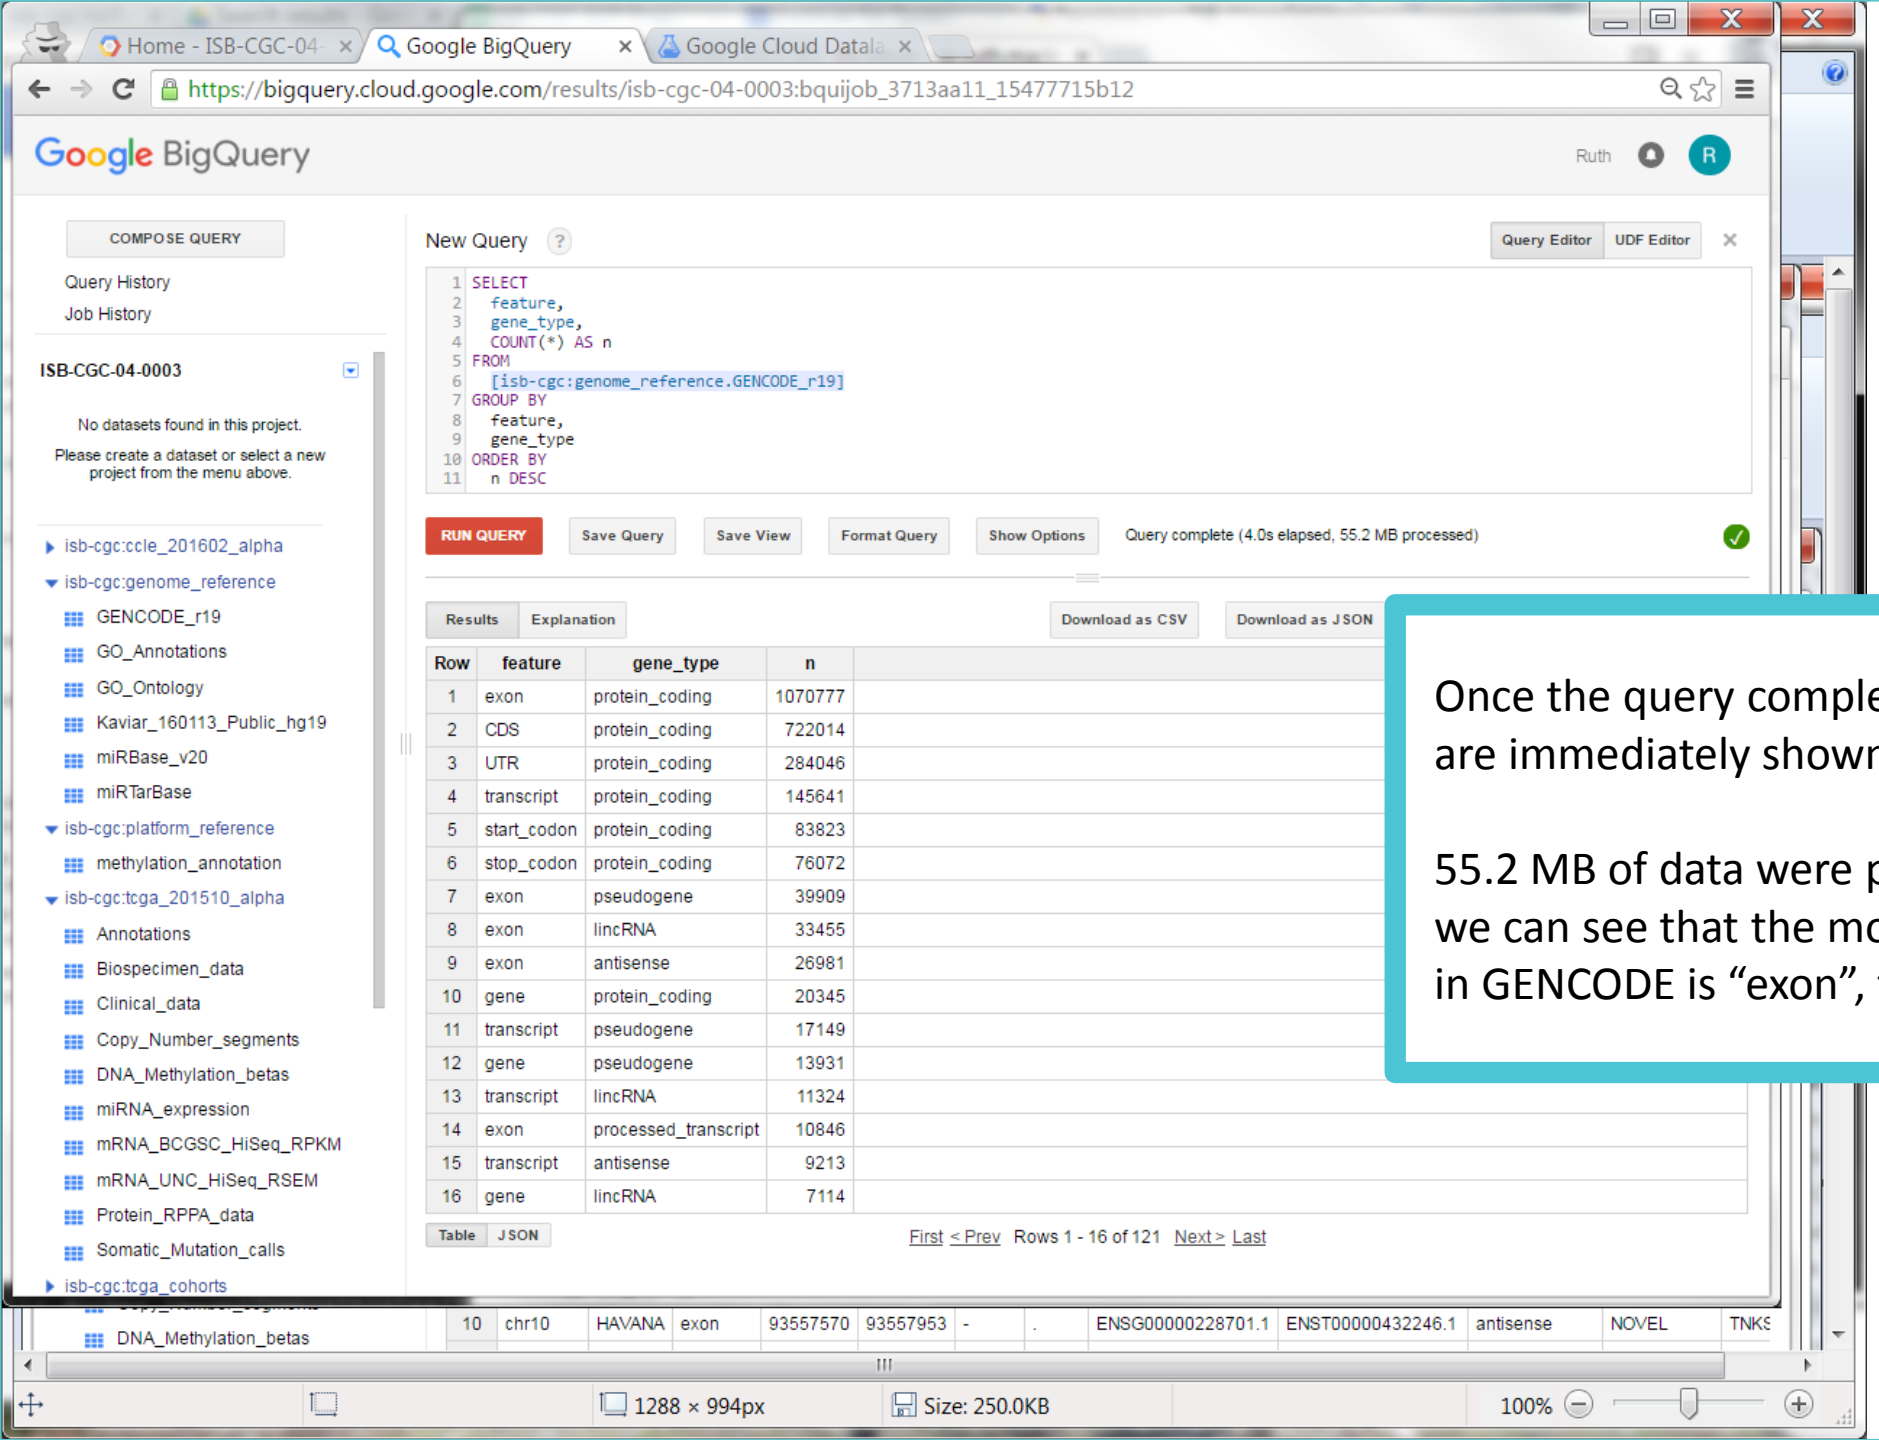

Once the query completes successfully, the results are immediately shown in the lower pane.

55.2 MB of data were processed in 4 seconds, and we can see that the most common type of feature in GENCODE is “exon”, followed by “CDS” *etc*

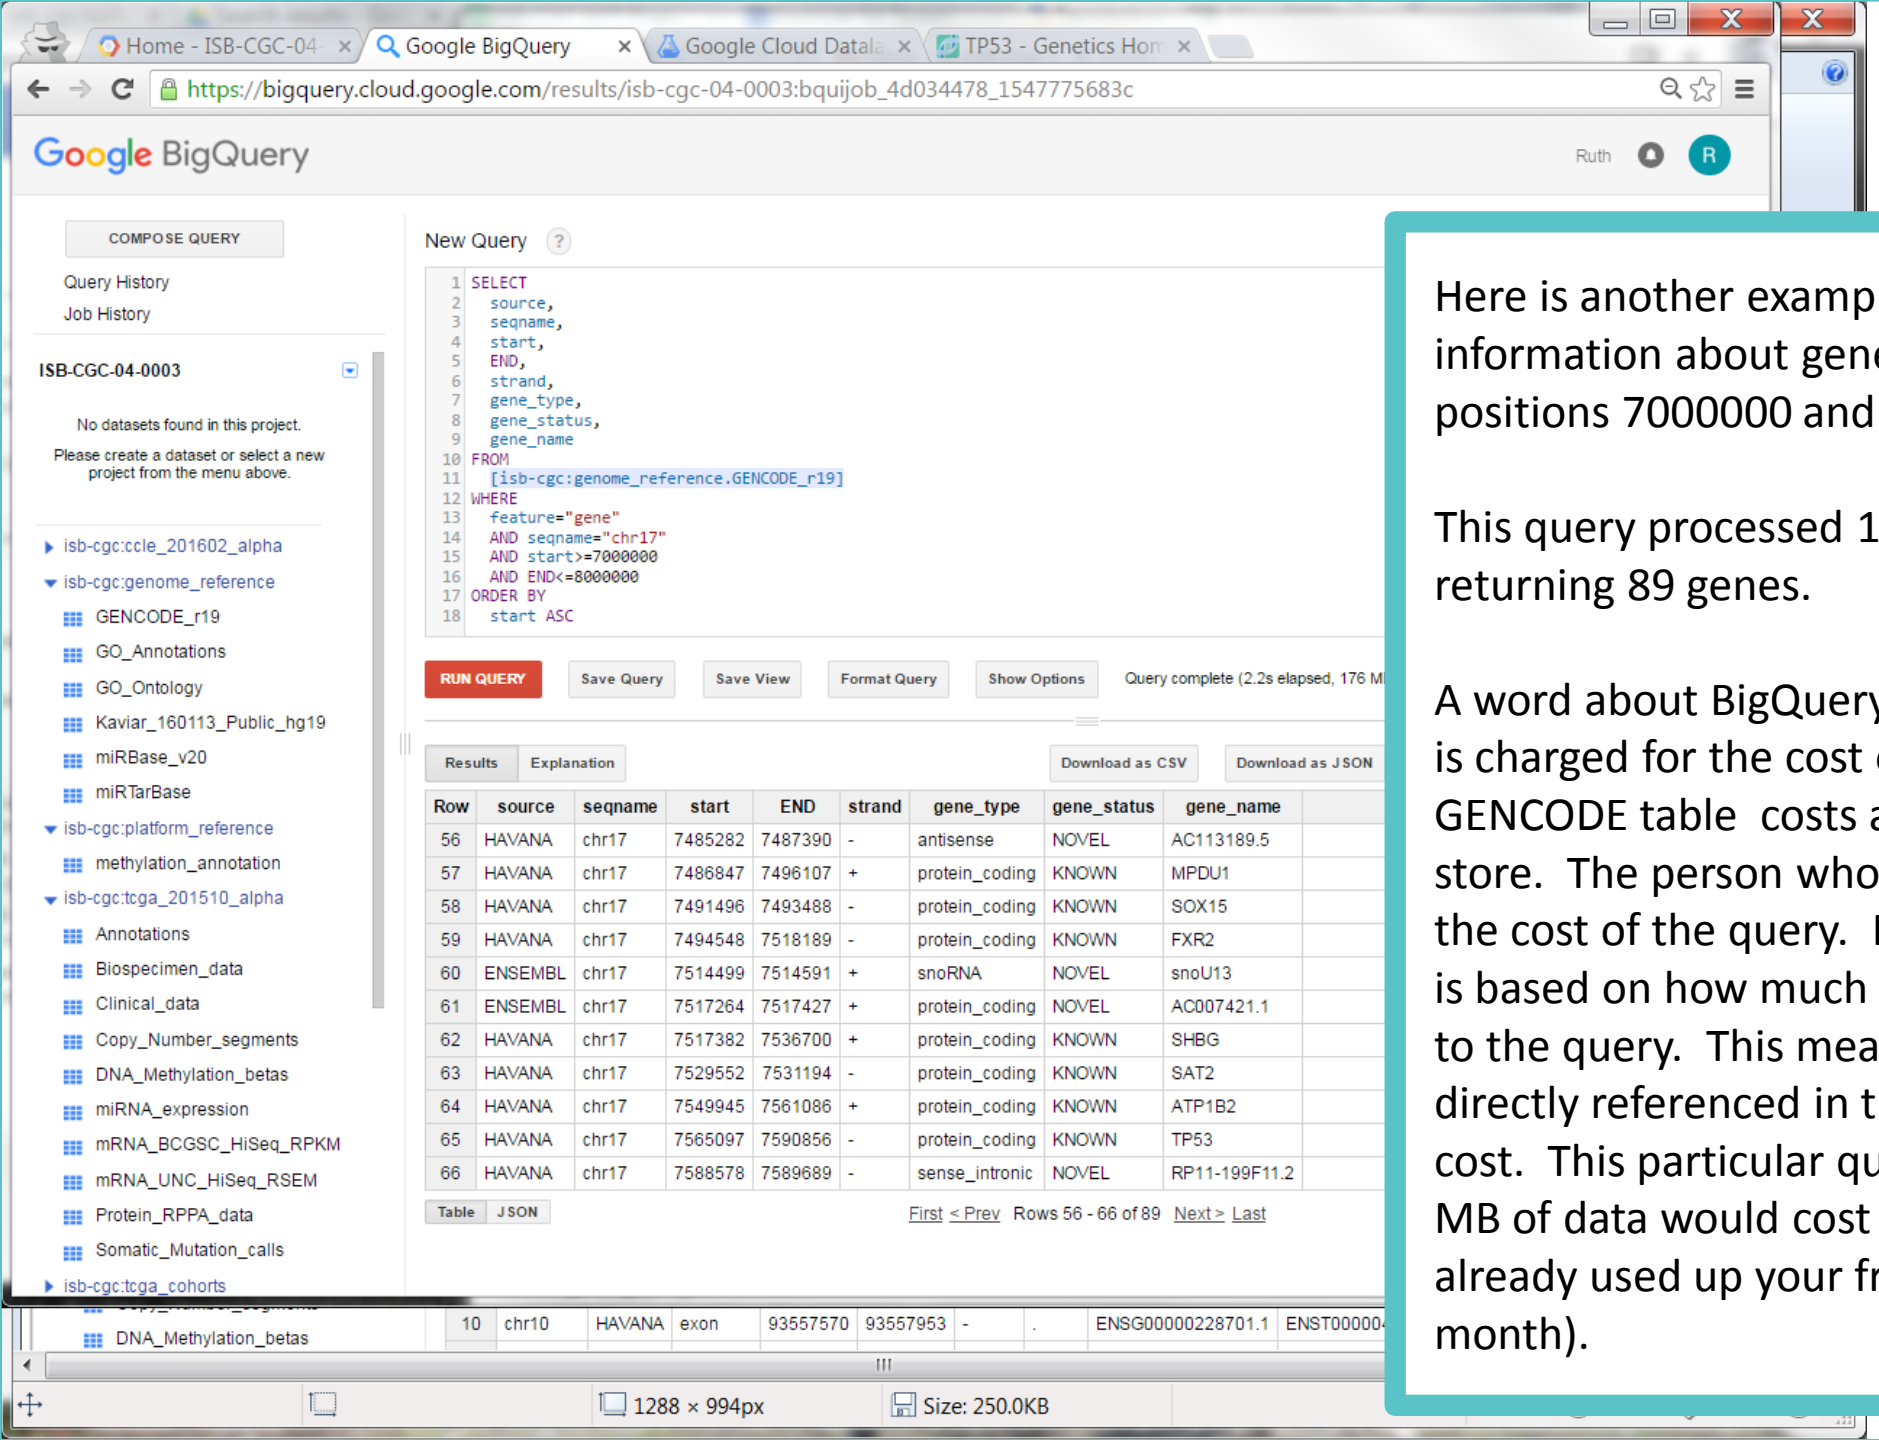

Here is another example query, which asks for information about genes on chr17 between positions 7000000 and 8000000.

This query processed 176 MB in just 2.2 seconds, returning 89 genes.

A word about BigQuery costs. The owner of a table is charged for the cost of the storage, and this GENCODE table costs about 7 cents per year to store. The person who runs a query gets charged the cost of the query. For most queries, this charge is based on how much data is "scanned" to respond to the query. This means only columns that are directly referenced in the query count towards the cost. This particular query, which processed 176 MB of data would cost less than one cent (if you've already used up your free $5 worth of queries this month).

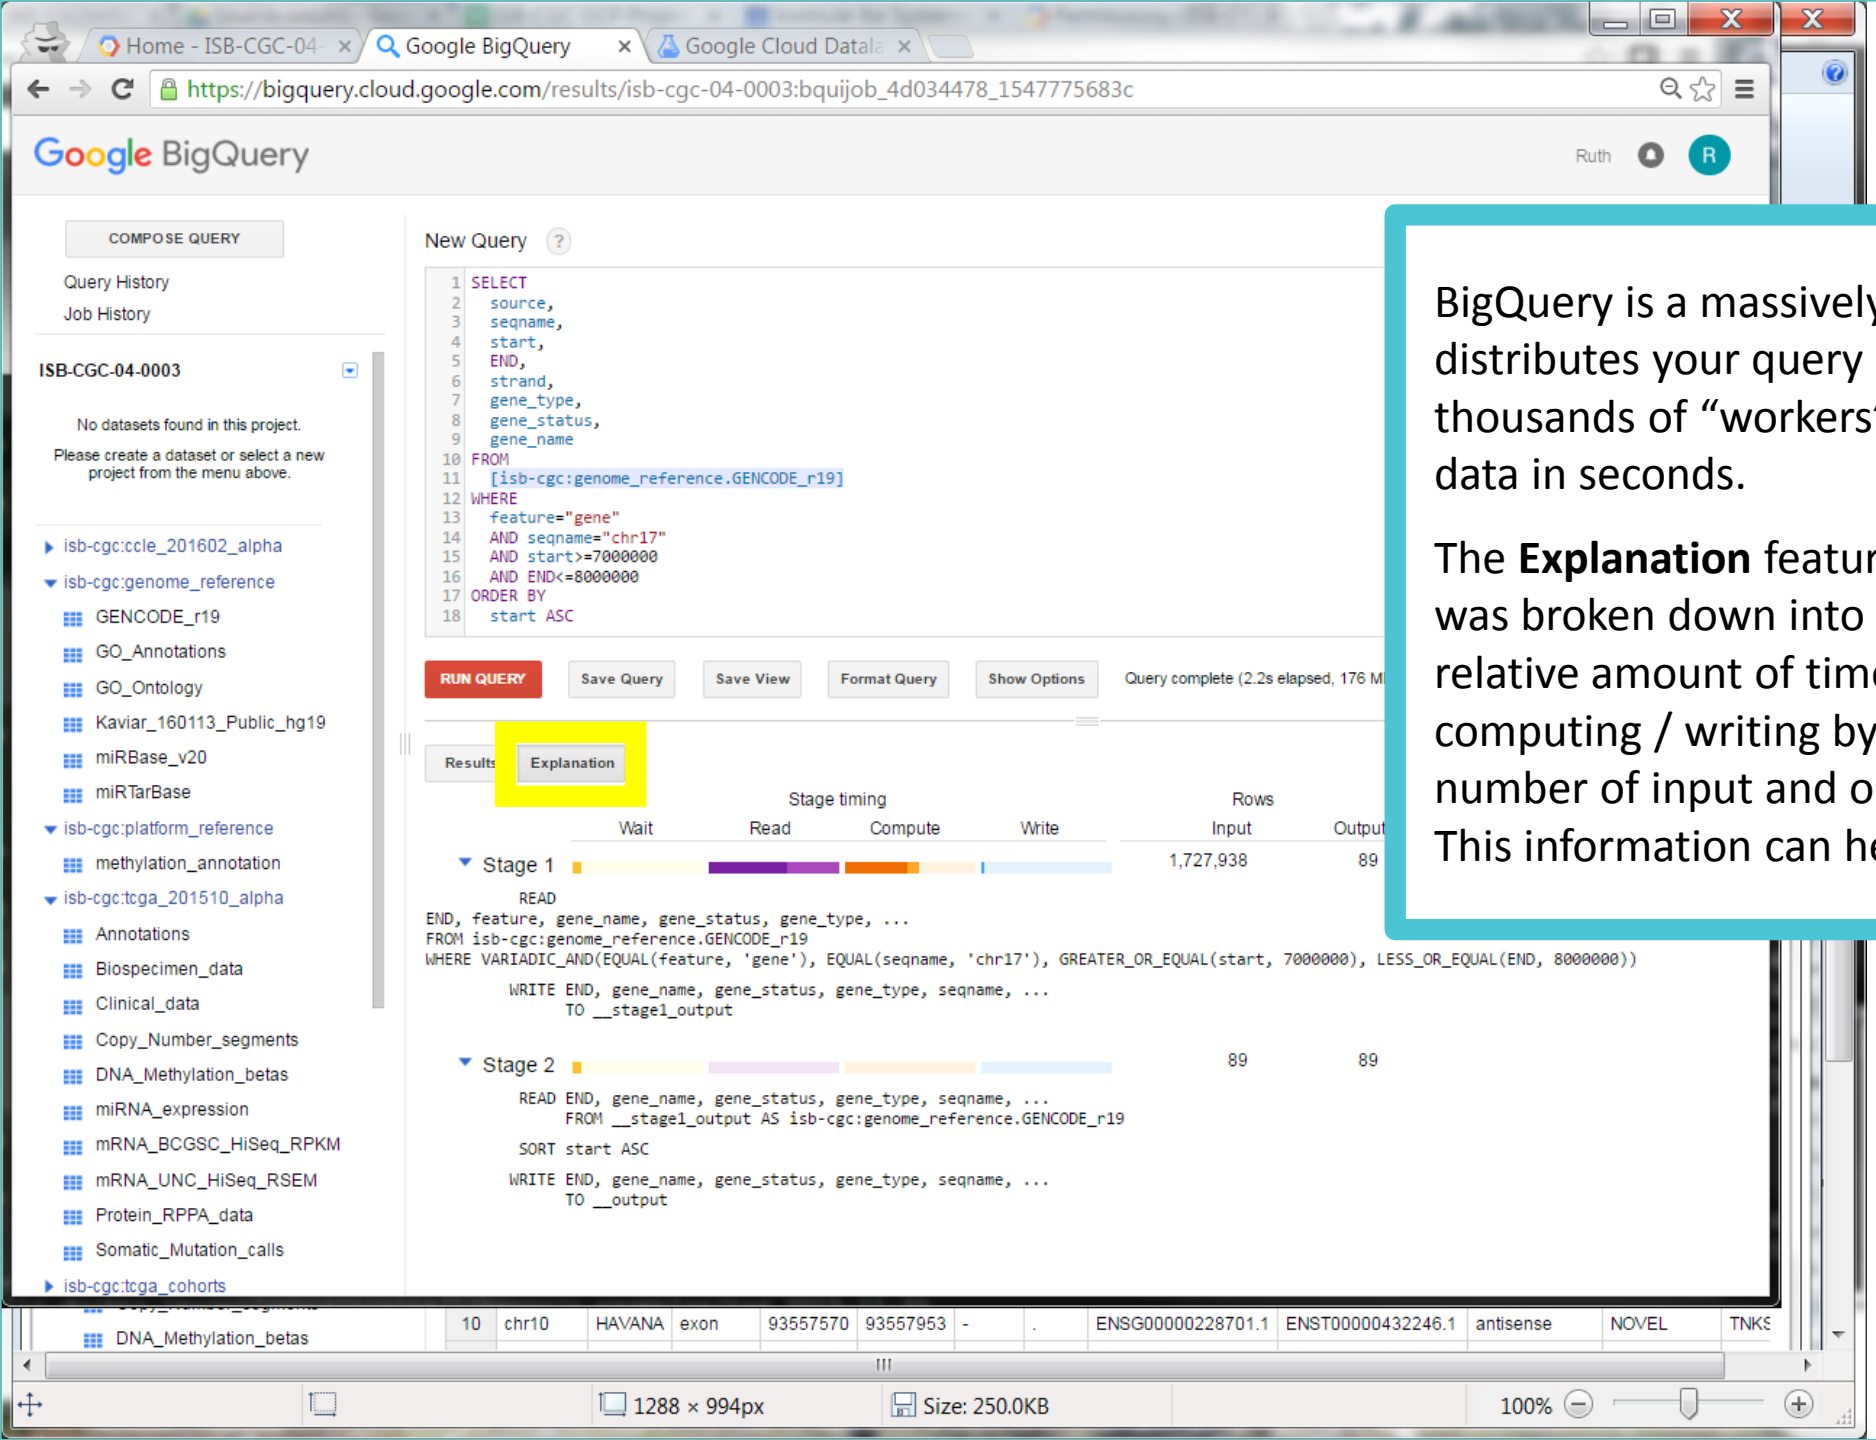

BigQuery is a massively parallel engine which distributes your query across hundreds or thousands of “workers” and can scan terabytes of data in seconds.

The **Explanation** feature shows you how your query was broken down into a series of stages, the relative amount of time spent waiting / reading / computing / writing by the “workers”, and the number of input and output rows at each stage. This information can help you optimize your query.

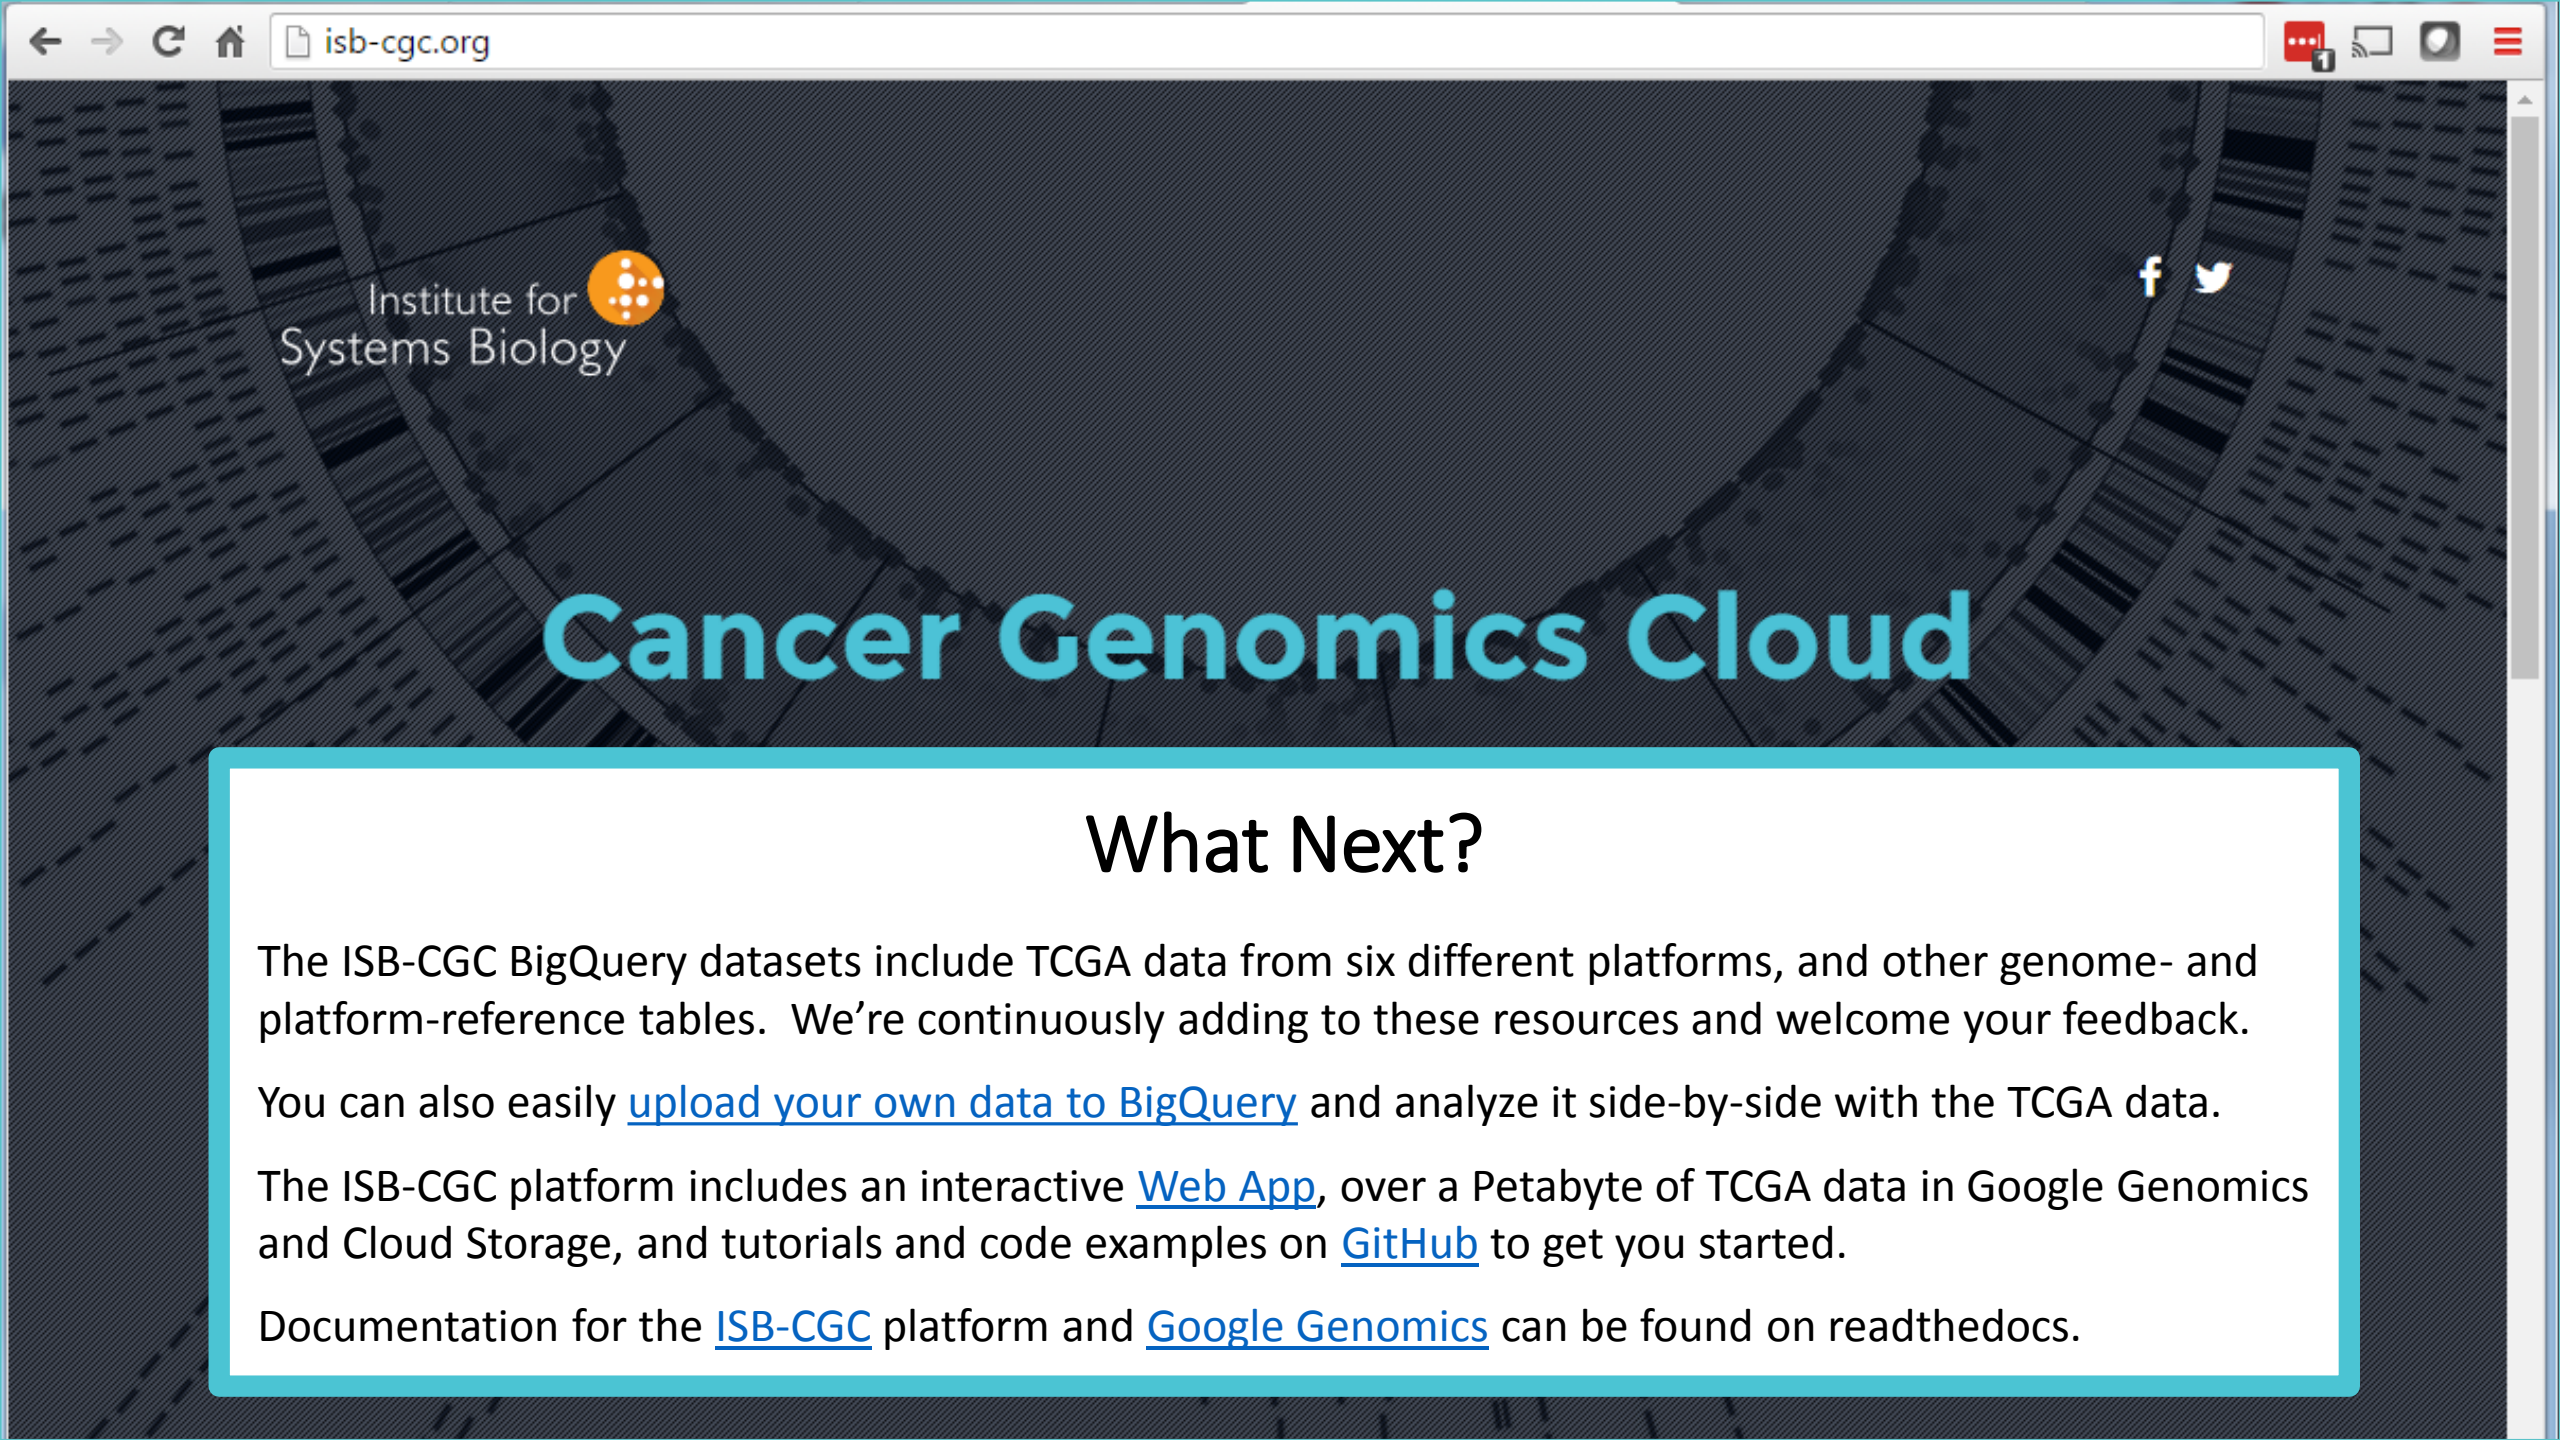
isb-cgc.org
Institute for Systems Biology
Cancer Genomics Cloud
What Next?
The ISB-CGC BigQuery datasets include TCGA data from six different platforms, and other genome- and platform-reference tables. We're continuously adding to these resources and welcome your feedback.
You can also easily upload your own data to BigQuery and analyze it side-by-side with the TCGA data.
The ISB-CGC platform includes an interactive Web App, over a Petabyte of TCGA data in Google Genomics and Cloud Storage, and tutorials and code examples on GitHub to get you started.
Documentation for the ISB-CGC platform and Google Genomics can be found on readthedocs.